

- [1 Introduction](#)
- [2 Installation](#)
- [3 Creating Users and Connections](#)
- [4 Managing Databases](#)
- [5 Creating, Changing and Removing Tables](#)
- [6 Managing Table Data](#)
- [7 Writing SQL Queries](#)
- [8 Visualizing Query Plans](#)
- [9 Visualizing Data](#)
- [10 Managing other PostgreSQL Elements](#)
- [11 Additional Features](#)
- [12 OmniDB Config Tool](#)
- [13 Writing and Debugging PL/pgSQL Functions](#)
- [14 Monitoring Dashboard](#)

1 Introduction

OmniDB is an open source browser-based app designed to access and manage many different Database Management systems, e.g. PostgreSQL, Oracle and MySQL. OmniDB can run either as an App or via Browser, combining the flexibility needed for various access paths with a design that puts security first. OmniDB is actively developed, automatically tested on a variety of databases and browsers and comes with full documentation.

Since early development, OmniDB was designed as an browser-based app. Consequently, it runs in any browser, from any operational system. It can be accessed by several computers and multiple users, each one of them with his/her own group of connections. It also can be hosted in any operational system, without the need of install any dependencies. We will see further details on installation in the next chapters.

OmniDB's main objective is to offer an unified workspace with all functionalities needed to manipulate different DMBS. DBMS specific tools aren't required: in OmniDB, the context switch between different DBMS is done with a simple connection switch, without leaving the same page. The end-user's sensation is that there is no difference when he/she manipulates different DBMS, it just feels like different connections.

Despite this, OmniDB is built with simplicity in mind, designed to be a fast and lightweight browser-based application. OmniDB is also powered by the WebSocket technology, allowing the user to execute multiple queries and procedures in multiple databases in multiple hosts in background.

OmniDB is also secure. All OmniDB user data are stored encrypted, and no database password is stored at all. When the user first connects to a database, OmniDB asks for the password. This password is encrypted and stored in memory for a specific amount of time. When this time expires, OmniDB asks the password again. This ensures maximum security for the database OmniDB is connecting to.

1.0.0.1 History

OmniDB's creators, Rafael Thofehrn Castro and William Ivanski, worked in a company where they needed to deal with several different databases from customers on a daily basis. These databases were from different DBMS technologies, and so they needed to keep switching between database management tools (typically one for each DBMS). As they were not keen of the existing unified database management tools (that could manage different DBMS), they came up with OmniDB's main idea.

OmniDB's first version was presented as an undergrad final project in the Computer Science Course from the Federal University of Paraná, in Brazil. The objective was to trace a common line between popular DBMS, and to study deeply their *metadata*. The result was a tool written in ASP.NET/C# capable of connecting and identifying the main structures (tables, keys, indexes and constraints), in a generic way, from several DBMS:

- Firebird
- MariaDB / MySQL
- Oracle
- PostgreSQL
- SQLite
- Microsoft SQL Server

OmniDB's first version also allowed the conversion between all DBMSs supported by the tool. This feature was developed to be user friendly, requiring just a few steps: the user needs to select a source connection, the structures that will be converted (just tables and all their structures, along with their data) and the target connection.

2 Installation

OmniDB provides 2 kinds of packages to fit every user needs:

- **OmniDB Application:** Runs a web server on a random port behind, and provides a simplified web browser window to use OmniDB interface without any additional setup. Just feels like a desktop application.
- **OmniDB Server:** Runs a web server on a random port. User needs to connect to it through a web browser. Provides user management, ideal to be hosted on a server on users' networks.

Both application and server can be installed on the same machine.

2.0.0.1 OmniDB Application

In order to run OmniDB app, you don't need to install any additional piece of software. Just head to omnidb.org and download the latest package for your specific operating system and architecture:

- Linux 32 bits / 64 bits
 - DEB installer
 - RPM installer
- Windows 32 bits / 64 bits
 - EXE installer
- Mac OSX
 - DMG installer

Use the specific installer for your Operational System and it will be available through your desktop environment application menu or via command line with `omnidb-app`.



2.0.0.2 OmniDB Server

Like OmniDB app, OmniDB server doesn't require any additional piece of software and the same options for operating system and architecture are provided.

Use the specific installer for your Operational System and it will be available through command line with `omnidb-server`.

You will need administrator privileges:

```
user@machine:~$ sudo omnidb-server
Starting OmniDB websocket...
Checking port availability...
Starting websocket server at port 25482.
OmniDB successfully migrated user database from version 0.0.0 to version 2.4.0
Starting OmniDB server...
Checking port availability...
Starting server OmniDB 2.4.0 at 0.0.0.0:8000.
Open OmniDB in your favorite browser
Press Ctrl+C to exit
```

Note how OmniDB starts a *websocket server* in a random port (in the above example it started in port 25482) and a *web server* in port 8000. You can also specify websocket server port, web server port and listening address:

```
user@machine:~$ sudo omnidb-server -p 8080 -w 2000 -H 127.0.0.1
Starting OmniDB websocket...
Checking port availability...
Starting websocket server at port 2000.
OmniDB successfully migrated user database from version 0.0.0 to version 2.4.0
Starting OmniDB server...
Checking port availability...
Starting server OmniDB 2.4.0 at 127.0.0.1:8080.
Open OmniDB in your favorite browser
```

2.0.0.3 OmniDB User Database

Since version 2.4.0, upon initialization both server and app will create a file `~/.omnidb/omnidb-app/omnidb.db` (for OmniDB app) or `~/omnidb/omnidb-server/omnidb.db` (for OmniDB server) in the user home directory, if it does not exist. That can be confirmed by the message *OmniDB successfully migrated user database from version 0.0.0 to version 2.4.0* you saw above. This file is also called **user database** and contains user data. If it already exists, then OmniDB will check whether the version of the server matches the version of the user database:

```
user@machine:~$ sudo omnidb-server
Starting OmniDB websocket...
Checking port availability...
Starting websocket server at port 25482.
User database version 2.4.0 is already matching server version.
Starting OmniDB server...
Checking port availability...
Starting server OmniDB 2.4.0 at 0.0.0.0:8000.
Open OmniDB in your favorite browser
Press Ctrl+C to exit
```

Future releases of OmniDB will contain the **user database migration** SQL commands required to upgrade the user database, if necessary. This way user data is not lost by upgrading OmniDB. Imagine the following scenario: you use OmniDB 2.4.0 now and you decide to upgrade it to newest release 2.5.0, for example. After the upgrade, when you start OmniDB server, it will apply the changes version 2.5.0 requires. So you will see something like that:

```
user@machine:~$ sudo omnidb-server
Starting OmniDB websocket...
Checking port availability...
Starting websocket server at port 25482.
OmniDB successfully migrated user database from version 2.4.0 to version 2.5.0
```

```
Starting OmniDB server...
Checking port availability...
Starting server OmniDB 2.5.0 at 0.0.0.0:8000.
Open OmniDB in your favorite browser
Press Ctrl+C to exit
```

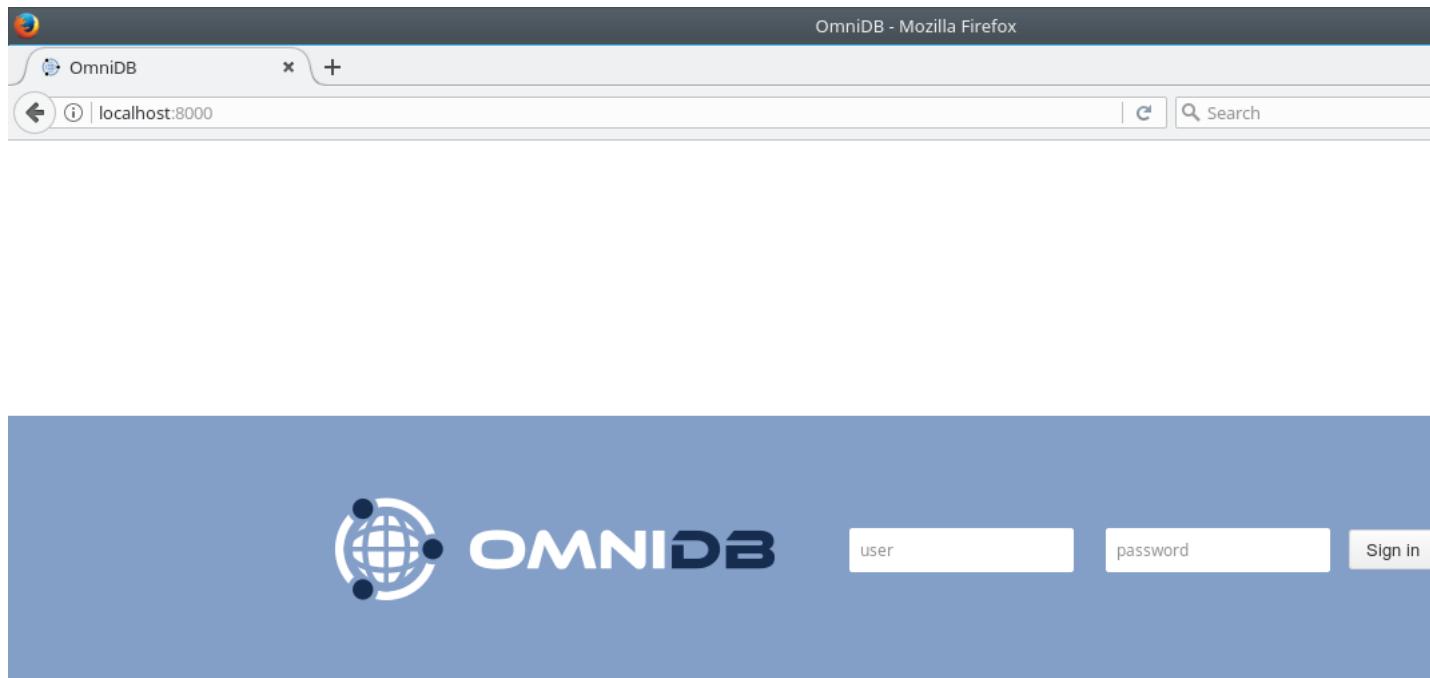
2.0.0.4 OmniDB configuration file

Starting on version 2.1.0, OmniDB server comes with a configuration file `omnidb.conf` that enables the user to specify parameters such as port and listening address. Also, 2.1.0 enables us to start the server with SSL, this requires a certificate and is configured in the same configuration file.

Starting on version 2.4.0, this file is located in `~/.omnidb/omnidb-server/omnidb.conf` in the user home directory.

2.0.0.5 OmniDB in the browser

Now that the web server is running, you may access OmniDB browser-based app on your favorite browser. Type in address bar: `localhost:8000` and hit `Enter`. If everything went fine, you shall see a page like this:

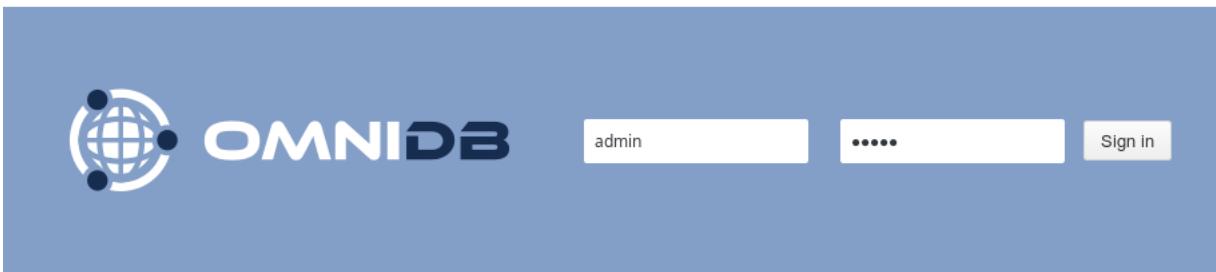


Now you know that OmniDB is running correctly. In the next chapters, we will see how to login for the first time, how to create an user and to utilize OmniDB.

3 Creating Users and Connections

3.0.0.1 Logging in as user *admin*

OmniDB comes only with the user *admin*. If you are using the server version, the first thing to do is sign in as *admin*, the default password is *admin*. You don't need to login in the app version.



The next window is the **Connections** window. We will talk about it later.

The image shows the OmniDB Connections window. At the top, there is a dark header bar with the OmniDB logo and the word "Workspace". Below the header is a toolbar with a "New Connection" button and several tabs: Technology, Server, Port, Database, User, Title, and Actions. The "User" tab is currently selected.

3.0.0.2 Creating another user

Click on the *Users* icon on the upper right corner. It will open a popup that allows the current OmniDB super user to create a new OmniDB user.

The image shows a modal dialog box titled "New User". It contains a table with four columns: Username, Password, Super User, and Actions. The "Username" column has "admin" in it. The "Password" column has a series of asterisks. The "Super User" column has a checked checkbox. The "Actions" column has a red "X" icon. In the top right corner of the dialog box is a blue circular close button with a white "X".

After clicking on the *Users* icon the tool inserts a new user called *user2* (if that is the first user after *admin*).

New User			
Username	Password	Super User	Actions
admin	*****	<input checked="" type="checkbox"/>	
user2	*****	<input type="checkbox"/>	

You will have to change the *username* and *password*. Check if you want this new user to be a *super user*. This user management window is only seen by super users. When you are done, click on the *Save Data* button inside the popup.

New User				Save Data
Username	Password	Super User	Actions	
admin	*****	<input checked="" type="checkbox"/>		
test	****	<input type="checkbox"/>		

You can create as many users as you want, edit existing users and also delete users by clicking on the red cross at the actions column. Now you can logout.

3.0.0.3 Signing in as the new user

Let us sign in as the user we just created.



And we can see the **Connections** window again. Note that now there is no *Users* icon, because the *test* user is not a super user.

OMNIDB		Workspace				
New Connection						
Technology	Server	Port	Database	User	Title	Actions

3.0.0.4 Creating connections

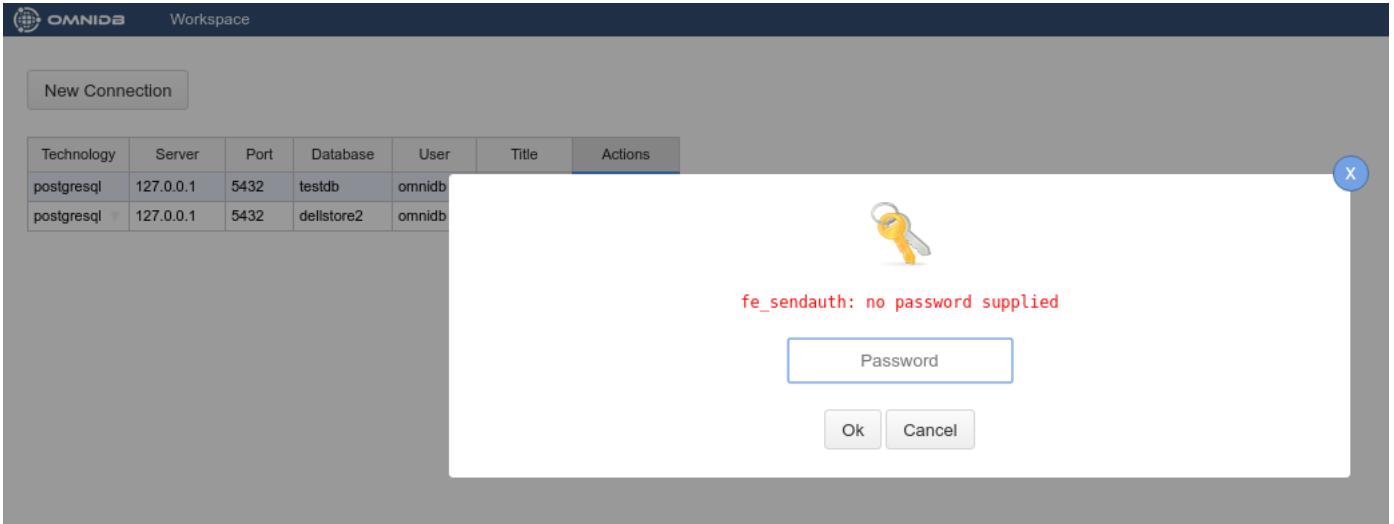
OmniDB C# version supported several DBMS. At the moment, OmniDB Python version, or OmniDB 2.0, supports only PostgreSQL. More DBMS support is being added as you read this.

We will now create two connections to PostgreSQL databases. To create the connections you have to click on the button *New Connection* and then choose the connection and fill all the fields for both connections, click on the *Save Data* button.

New Connection

Technology	Server	Port	Database	User	Title	Actions
postgresql	127.0.0.1	5432	testdb	omnidb	Test	  
postgresql	127.0.0.1	5432	dellstore2	omnidb	DellStore	  

For each connection there is an *Actions* column where you can delete, test and select them. Go ahead and test one of the connections.

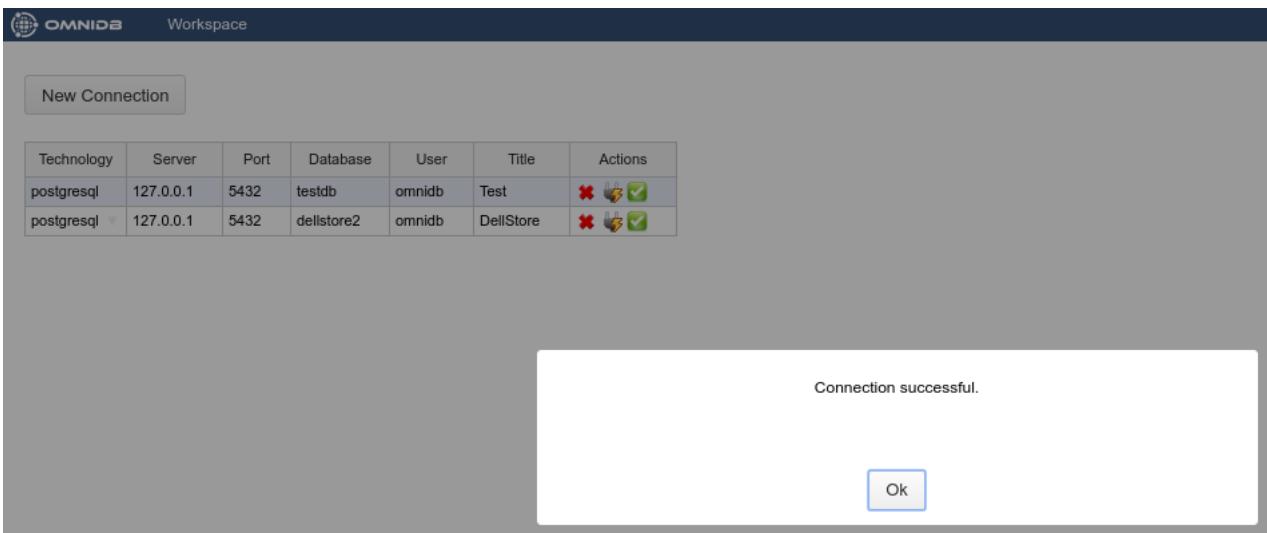


The screenshot shows the OmniDB interface with a 'New Connection' dialog open. The dialog contains a table with two rows of connection details. A modal window is overlaid on the main interface, prompting for a password. The modal has a yellow key icon at the top, followed by the error message 'fe_sendauth: no password supplied'. Below the message is a 'Password' input field, and at the bottom are 'Ok' and 'Cancel' buttons. A close button ('X') is located in the top right corner of the modal.

Technology	Server	Port	Database	User	Title	Actions
postgresql	127.0.0.1	5432	testdb	omnidb		
postgresql	127.0.0.1	5432	dellstore2	omnidb		

Notice a pop-up appears with the message *fe_sendauth: no password supplied*. This is happening because OmniDB does not store the database user password on disk. Not having any password at hand, OmniDB will try to connect without one, thus trying to take advantage of automatic authentication methods that might be in place: `trust` method, `.pgpass` file, and so on. As the database server replies with an error not allowing the user to connect, then OmniDB understands a password is required and asks it to the user. When the user types a password in this popup, the password is encrypted and stored in memory.

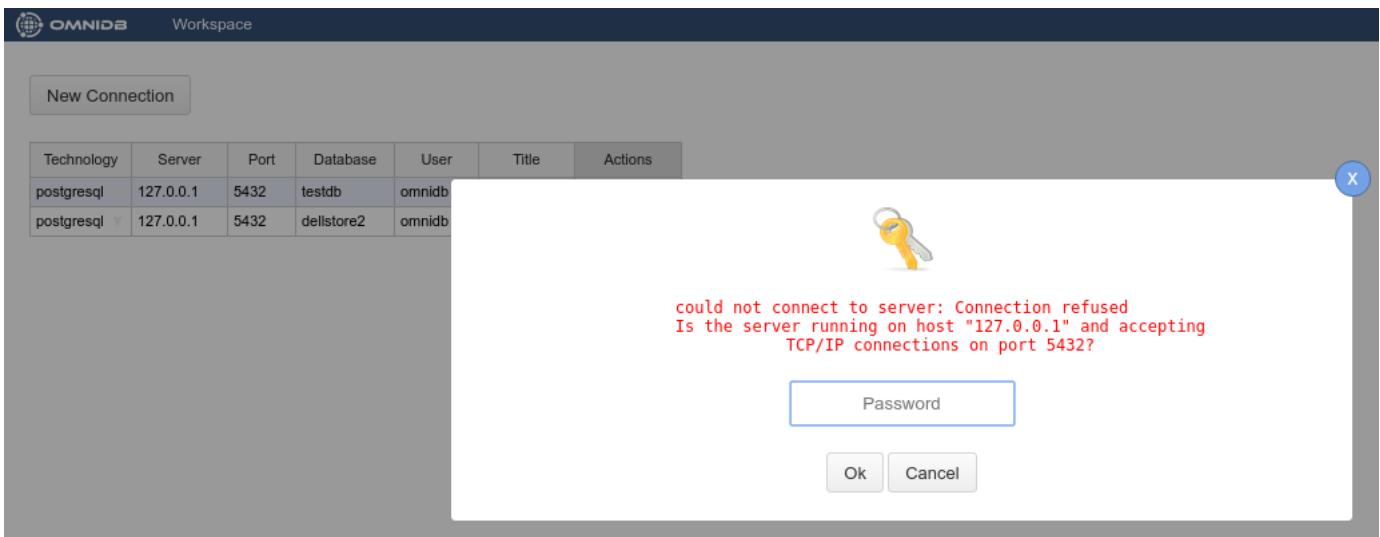
After you type the password and hit *Enter*, if the connection to the database is successful you will see a confirmation pop-up.



The screenshot shows the OmniDB interface with a 'New Connection' dialog open. The dialog contains a table with two rows of connection details. A modal window is overlaid on the main interface, confirming the connection. The modal displays the message 'Connection successful.' and has an 'Ok' button at the bottom.

Technology	Server	Port	Database	User	Title	Actions
postgresql	127.0.0.1	5432	testdb	omnidb	Test	  
postgresql	127.0.0.1	5432	dellstore2	omnidb	DellStore	  

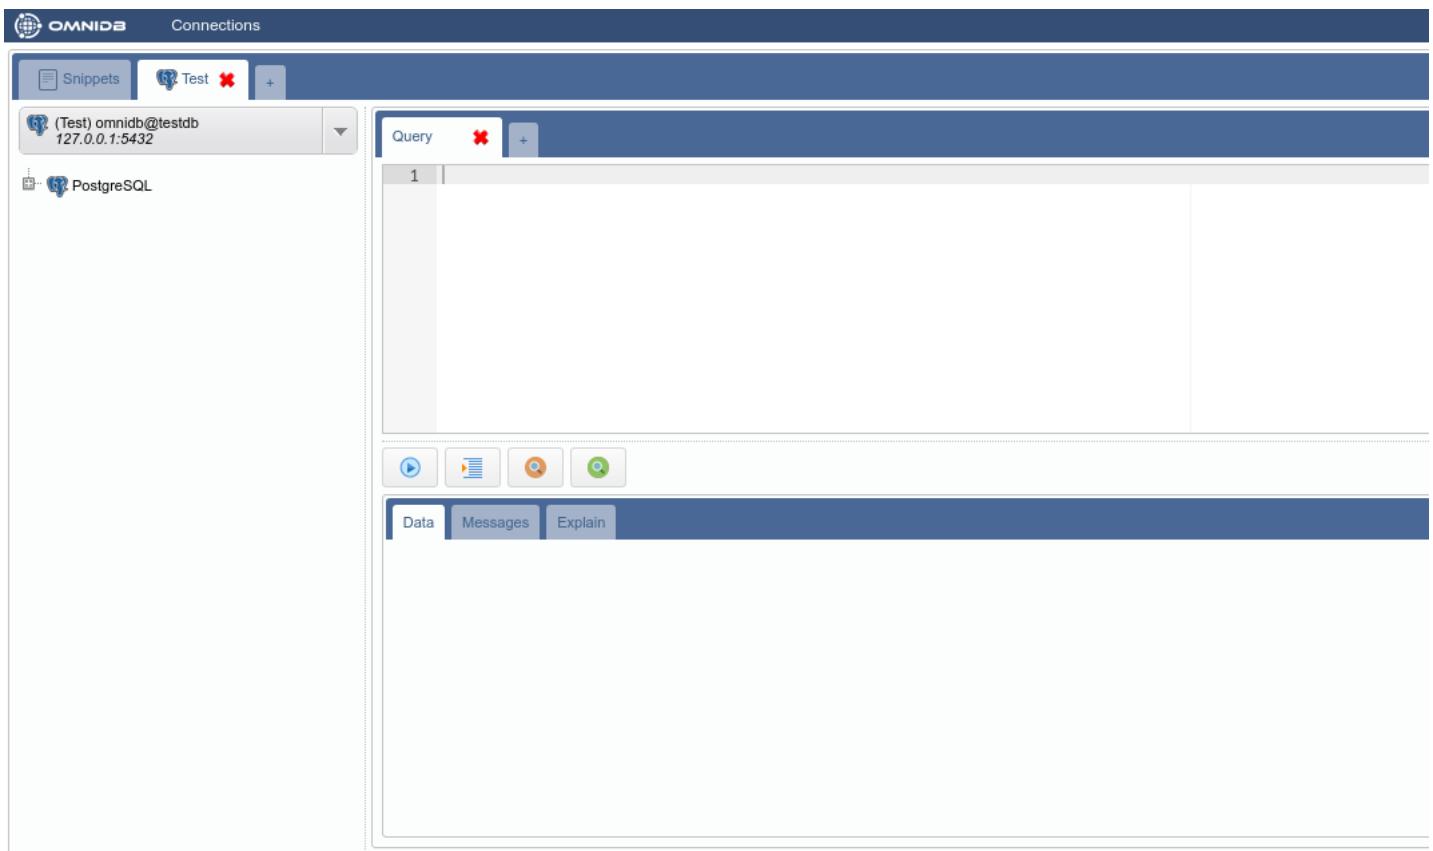
But, if you have trouble of any kind connecting to your PostgreSQL database, the same popup will remain showing the error OmniDB got.



Also, in the connections grid, if you click on the *Select Connection* action, OmniDB will open it in the **Workspace Window**.

4 Managing Databases

After creating at least one connection the user can enter the *Workspace*, either by clicking the *Workspace* tab or by clicking in the *Select Connection* action in the connections grid.



4.0.0.1 Sections of the *Workspace* window

This interface has several elements:



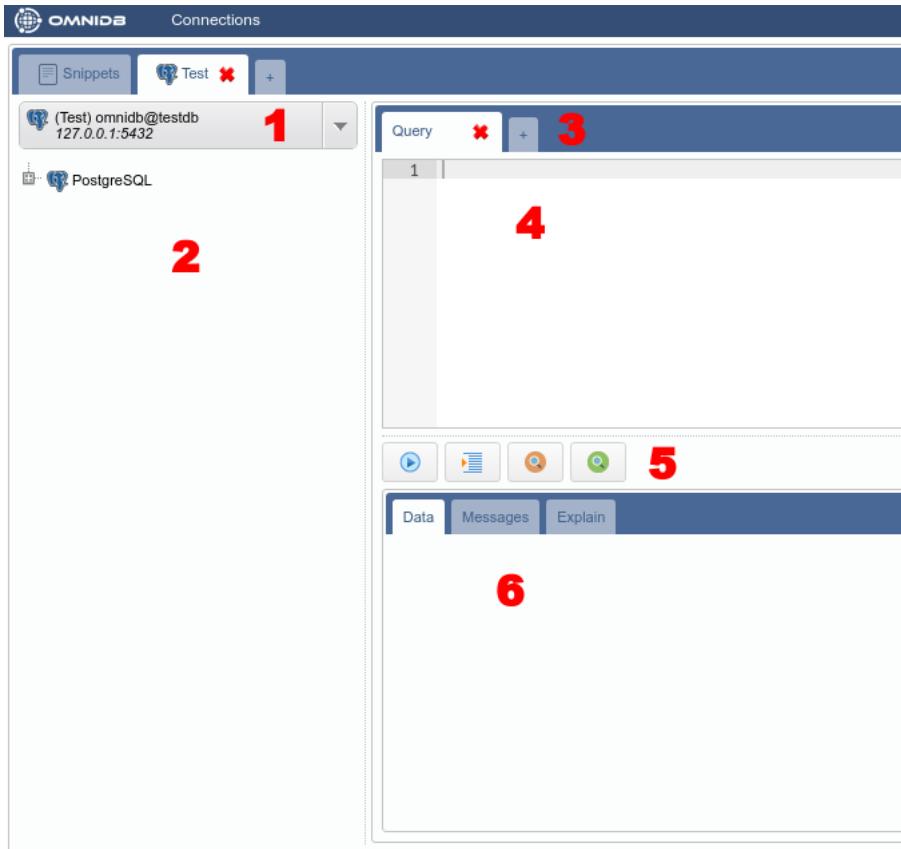
- **1)** **Links:** Enables the user to navigate to the *Connections* window and back to the *Workspace* window

- **2) Outer Tabs:** OmniDB lets you work with several databases at the same time. Each database will be accessible through an *outer tab*. Outer tabs also can host miscellaneous connection-independent features, like the *Snippets* feature
- **3) Options:** Shows the current user logged in, and also links for *user settings*, *query history*, *information* and *logout*.

4.0.0.2 Connection Outer Tab

So, the outer table named *Test* has this name because of the alias we put in the connection to the *testdb*. This tab is a *Connection Outer Tab*. Notice the little tab with a cross besides the *Test* outer tab. This allows you to create a new outer tab that will automatically be a *Connection Outer Tab*. However, the *Snippet Outer Tab* is fixed and will always be the first.

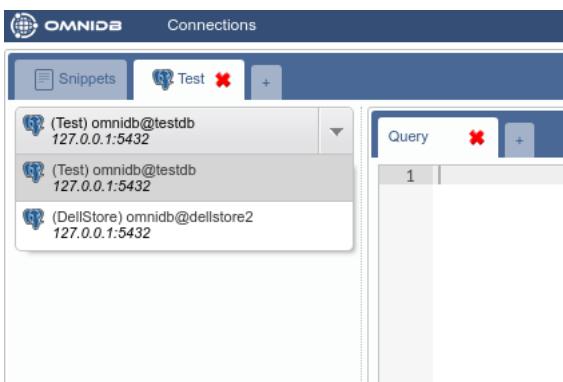
A new *Connection Outer Tab* will always automatically point to the first connection on your list of database connections. Or, if you clicked on the *Select Connection* action, it will point to the selected connection. Observe the elements inside of this tab:



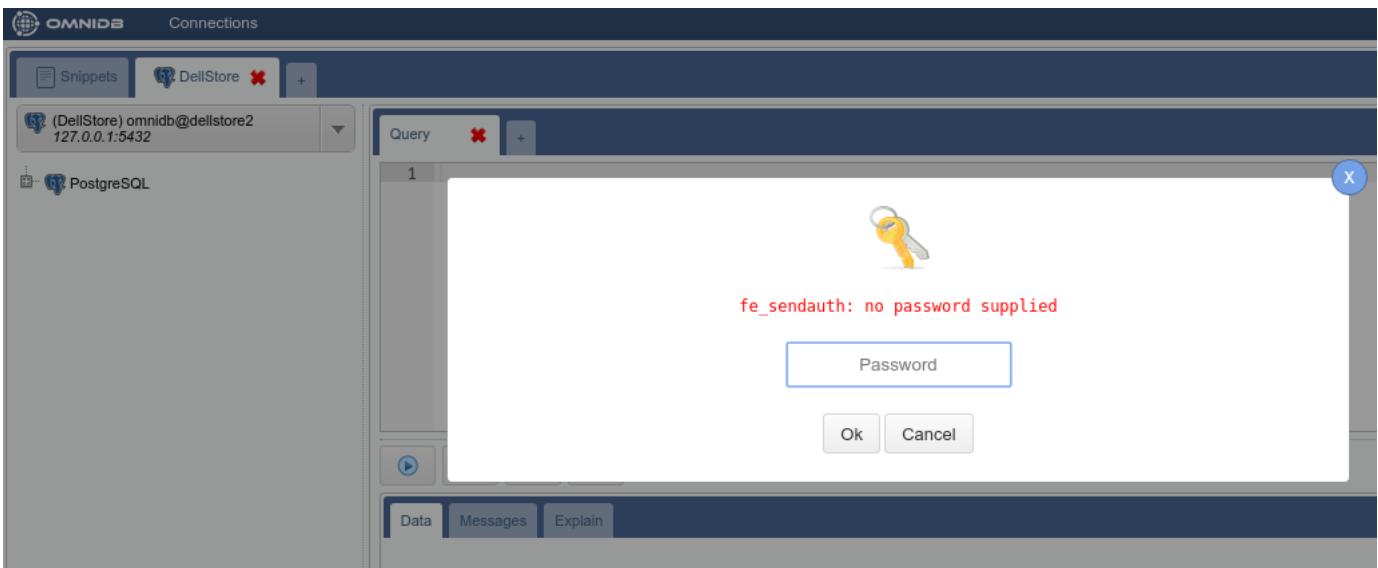
- **1) Connection Selector:** Shows all connections and lets the user select the current one
- **2) Tree of Structures:** Displays a hierarchical tree where you can navigate through the database elements
- **3) Inner Tabs:** Allows the user to execute actions in the current database. There are several kinds of inner tabs that will be a *Query Tab*, where you can write any kind of SQL statement
- **4) Inner Tab Content:** Can vary depending on the kind of inner tab. The figure shows a *Query Tab* and in this case the content will be an *SQL Editor*, with syntax highlight and autocomplete
- **5) Inner Tab Actions:** Can vary depending on the kind of inner tab. For a *Query Tab*, they are *Execute*, *Format*, *Explain* and *Explain Analyze*
- **6) Inner Tab Results:** A *Query Tab*, after you click in the *Execute Button* or type the execute shortcut (*Alt-Q*), will show a grid with the query results in the *Data* subtab. If the query calls a function that raises messages, those will be shown in the *Messages* subtab. If instead of *Execute* you clicked in *Explain* or *Explain Analyze*, the explain plan for the query will be shown in the *Explain* subtab.

4.0.0.3 Working with databases

Take a look at your connections selector. OmniDB always points to the first available connection but you can change it by clicking on the selector.

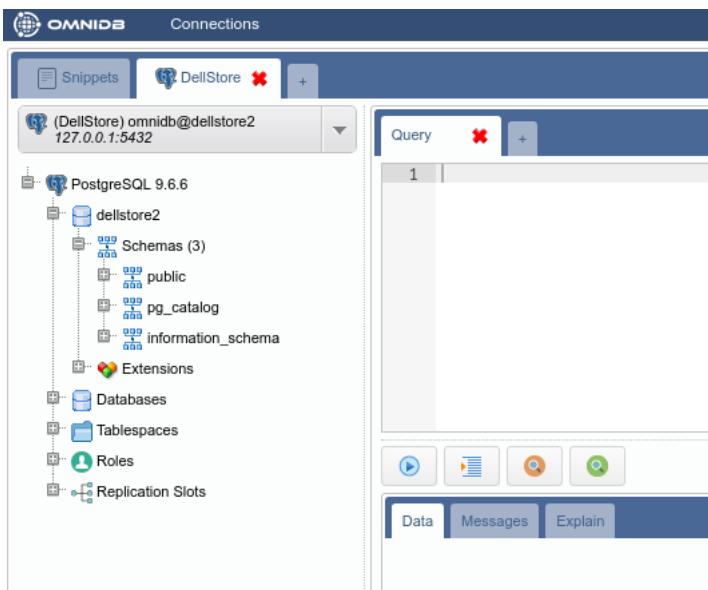


Select the *DellStore* connection. Now go to the tree right below the selector and click to expand the root node *PostgreSQL*.



Bear in mind that every 30 minutes you keep without performing actions on the database, will trigger a *Authentication* popup, meaning that the password that OmniDB has encrypted and stored in memory is now expired. As explained before, this is important for your database security. After you type the correct password, you will see the PostgreSQL node now shows the PostgreSQL version and also was expanded, showing the current database connection and also instance wide elements: *Databases*, *Tablespaces*, *Roles* and *Replication Slots*.

Go ahead and expand the *Schemas* node. You will see all schemas in the current database (in case of PostgreSQL, TOAST and temp schemas are not shown).



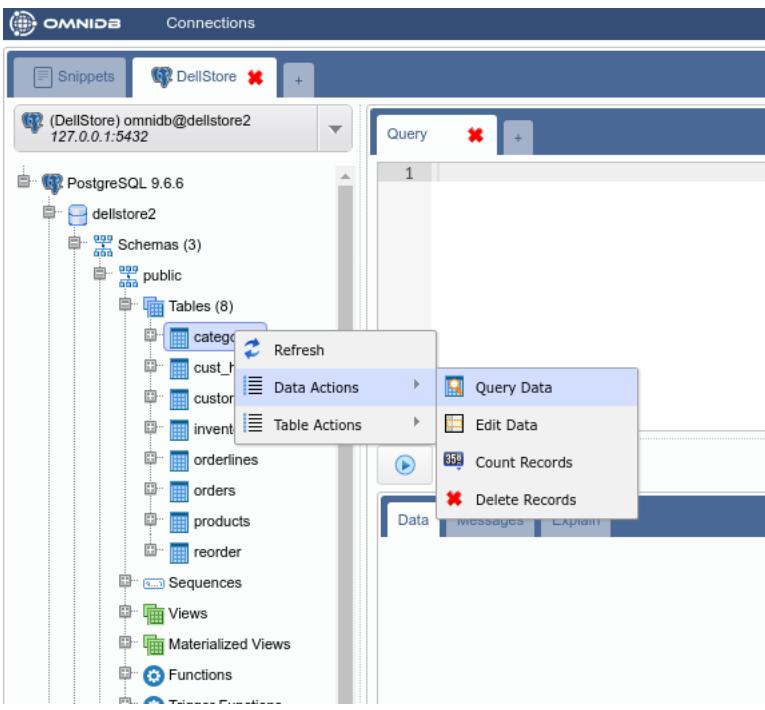
Now click to expand the schema `public`. You will see different kinds of elements contained in this schema.

The screenshot shows the OMNIDB interface for PostgreSQL 9.6.6. On the left, the schema browser displays the structure of the 'dellstore2' database, including the 'public' schema with its tables, sequences, views, materialized views, functions, trigger functions, and extensions. On the right, a query editor window titled 'Query' is open, showing a single row labeled '1'. Below the query editor are buttons for Data, Messages, and Explain.

Now click to expand the node *Tables*, and you will see all tables contained in the schema `public`. Expand any table and you will see its columns, primary key, foreign keys, constraints, indexes, rules, triggers and partitions.

This screenshot shows the same interface after expanding the 'Tables' node under the 'public' schema. The 'categories' table is now expanded, revealing its two columns: 'category' and 'categoryname'. The 'category' column is identified as a primary key ('Primary Key (1)'). Other table details like foreign keys, unique constraints, checks, excludes, indexes, rules, triggers, and partitions are also visible.

In order to view records inside a table, right click it and choose *Data Actions > Query Data*.



Notice that OmniDB opens a new SQL editor with a simple query to list table records. The records are displayed in a grid right below the editor. This grid can be controlled with keyboard as if you were using a spreadsheet manager. You can also copy data from single cells or block of cells (that can be selected with the keyboard or mouse) and paste on any spreadsheet manager.

	category	categoryname
1	1	Action
2	2	Animation
3	3	Children
4	4	Classics
5	5	Comedy
6	6	Documentary
7	7	Drama
8	8	Family
9	9	Foreign
10	10	Games
11	11	Horror
12	12	Music
13	13	New
14	14	Sci-Fi

You can edit the query on the SQL editor, writing simple or more complex queries. To execute, click on the action button or hit the keystroke **Ctrl-Q**. If the results exceed 50 registers, then extra buttons *Fetch More* and *Fetch All* will appear. More details in the next chapters.

4.0.0.4 Working with multiple tabs inside the same connection

Inside a single connection, you can create several inner tabs by clicking on the last little tab with a cross. Each new inner tab will be a *Query Tab*.

The screenshot shows the OmniDB interface with a connection to 'DellStore' at '127.0.0.1:5432'. The left sidebar displays the database schema for 'dellstore2', including 'Schemas (3)' and 'Tables (8)'. The 'categories' table is selected. The main area shows five tabs at the top: 'Query', 'categories', 'Query', 'Query', and 'Query'. The first tab contains the number '1'. Below the tabs is a toolbar with icons for play, stop, refresh, and search. At the bottom is a navigation bar with 'Data', 'Messages', and 'Explain' tabs.

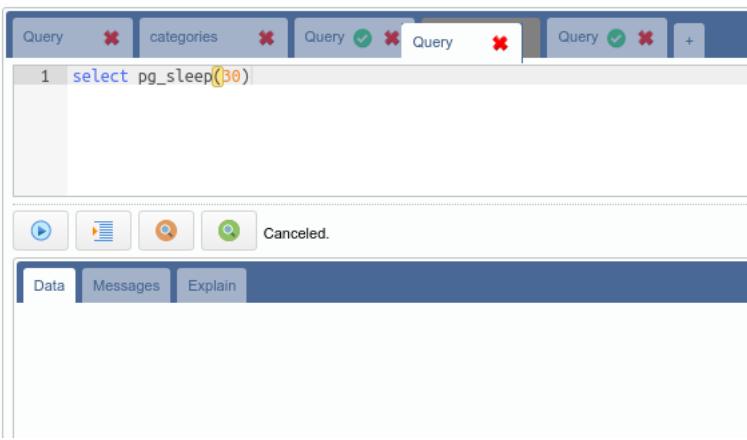
On OmniDB, you can execute several SQL statements and procedures in parallel. When it is executing, an icon will be shown in the tab to indicate its current state. If some process is finished and it is not in the current tab, that tab will show a green icon indicating the routine being executed there is now finished.

This screenshot shows the same OmniDB interface as above, but the 'categories' tab is now active. It contains the SQL statement 'select pg_sleep(30)'. A red button labeled 'Cancel' is visible. To its right, a message says 'Start time: 11/30/2017 16:39:46' and 'Running...'. The other tabs ('Query', 'Query', 'Query') are still present but inactive.

By clicking in the *Cancel* button, you can cancel a process running inside the database.

This screenshot shows the OmniDB interface after a cancellation. The 'categories' tab is active and shows the message 'Canceled.' in its content area. The other tabs ('Query', 'Query', 'Query') are still present but inactive.

You can also drag and drop a tab to change its order. This works with both inner and outer tabs.



Additionally, you can use keyboard shortcuts to manage inner tabs (SQL Query) and outer tabs (Connection):

- **Ctrl-Insert:** Insert a new inner tab
- **Ctrl-Delete:** Removes an inner tab
- **Ctrl-<:** Change focus to inner tab at left
- **Ctrl->:** Change focus to inner tab at right
- **Ctrl-Shift-Insert:** Insert a new outer tab
- **Ctrl-Shift-Delete:** Removes an outer tab
- **Ctrl-Shift-<:** Change focus to outer tab at left
- **Ctrl-Shift->:** Change focus to outer tab at right

Starting from OmniDB version 2.3.0, all SQL Query tabs are automatically saved whenever you execute them. Even if you close OmniDB window or browser tab, they are already stored in OmniDB *User Database*. They will be automatically restored when you open OmniDB again (if you are using app), open it in another browser window (if you are using server), or even if you clicked in the *Connections* window or logged out. Removing an outer tab or inner tab by the interface makes it permanently deleted, so it will not be restored.

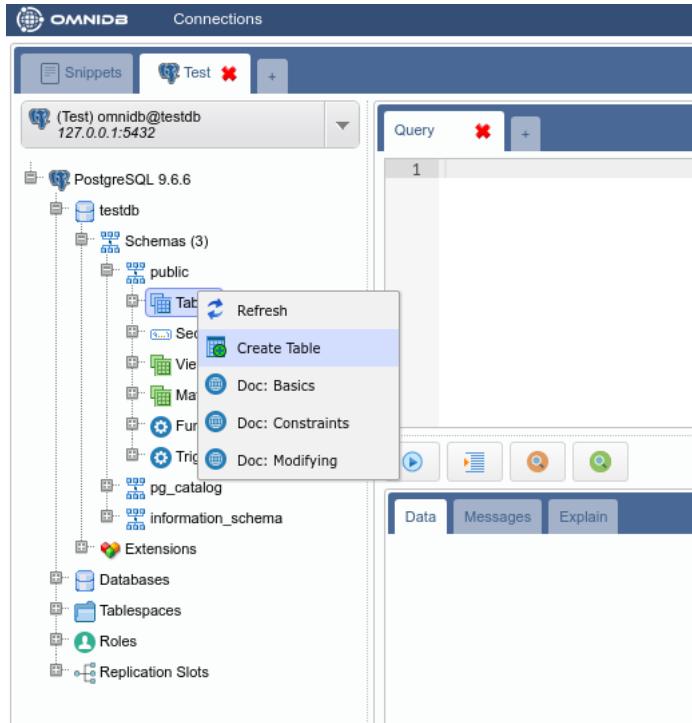
5 Creating, Changing and Removing Tables

5.0.0.1 Creating tables

OmniDB has a table creation interface that lets you configure columns, constraints and indexes. A couple of observations should be mentioned:

- Most DBMS automatically create indexes when primary keys and unique constraints are created. Because of that, the indexes tab is only available after creating the table.
- Each DBMS has its unique characteristics and limitations regarding table creation and the OmniDB interface reflects these limitations. For instance, SQLite does not allow us to change existing columns and constraints. Because of that, the interface lets us change only table name and add new columns when dealing with SQLite databases (it is still not the case in OmniDB Python version, as it currently supports only PostgreSQL databases).

We will create example tables (*Customer* and *Address*) in the *testdb* database we connected to earlier. Right click on the **Tables** node and select the **Create Table** action:



We will create the table *Customer* with a primary key that will be referenced by the table *Address*:

Query X New Table X +

Table Name: Customer Save Changes

Columns Constraints Indexes

	Column Name	Data Type	Nullable	
1	cust_id	serial	NO	X
2	cust_name	varchar(100)	NO	X
3				

Query X New Table X +

Table Name: Customer Save Changes

Columns Constraints Indexes

New Constraint

Constraint Name	Type	Columns	Referenced Table	Referenced Columns	Delete Rule	Update Rule	
cust_pk	Primary Key	cust_id					X

Click on the *Save Changes* button. Right-click the *Tables* tree node and click *Refresh*. Note how the table appears in the *Tables* tree node:

OMNIDB Connections

(Test) omnidb@testdb 127.0.0.1:5432

PostgreSQL 9.6.6

- testdb
 - Schemas (3)
 - public
 - Tables (1)
 - customer
 - Columns (2)
 - Primary Key (1)
 - cust_pk
 - cust_id
 - Foreign Keys
 - Uniques
 - Checks
 - Excludes
 - Indexes (1)
 - cust_pk (Unique)
 - cust_id
 - Rules
 - Triggers
 - Partitions

Now create the table *Address* with a primary key and a foreign key:

Query public.customer New Table

Table Name: Address Save Changes

Columns Constraints Indexes

	Column Name	Data Type	Nullable	
1	add_id	serial	NO	X
2	add_street	varchar(200)	NO	X
3	add_number	integer	YES	X
4	cust_id	integer	NO	X
5				

Query public.customer New Table

Table Name: Address Save Changes

Columns Constraints Indexes

New Constraint

Constraint Name	Type	Columns	Referenced Table	Referenced Columns	Delete Rule	Update Rule	
add_pk	Primary Key	<input type="checkbox"/> add_id					X
add_fk1	Foreign Key	<input type="checkbox"/> cust_id	public.customer	cust_id	CASCADE	CASCADE	X

Don't forget to click on the *Save Changes* button when done. At this point we have two tables in schema `public`. The schema structure can be seen with the graph feature by right clicking on the schema `public` node of the tree and selecting *Render Graph > Simple Graph*:

OMNIDB Connections

Snippets Test

(Test) omnidb@testdb 127.0.0.1:5432

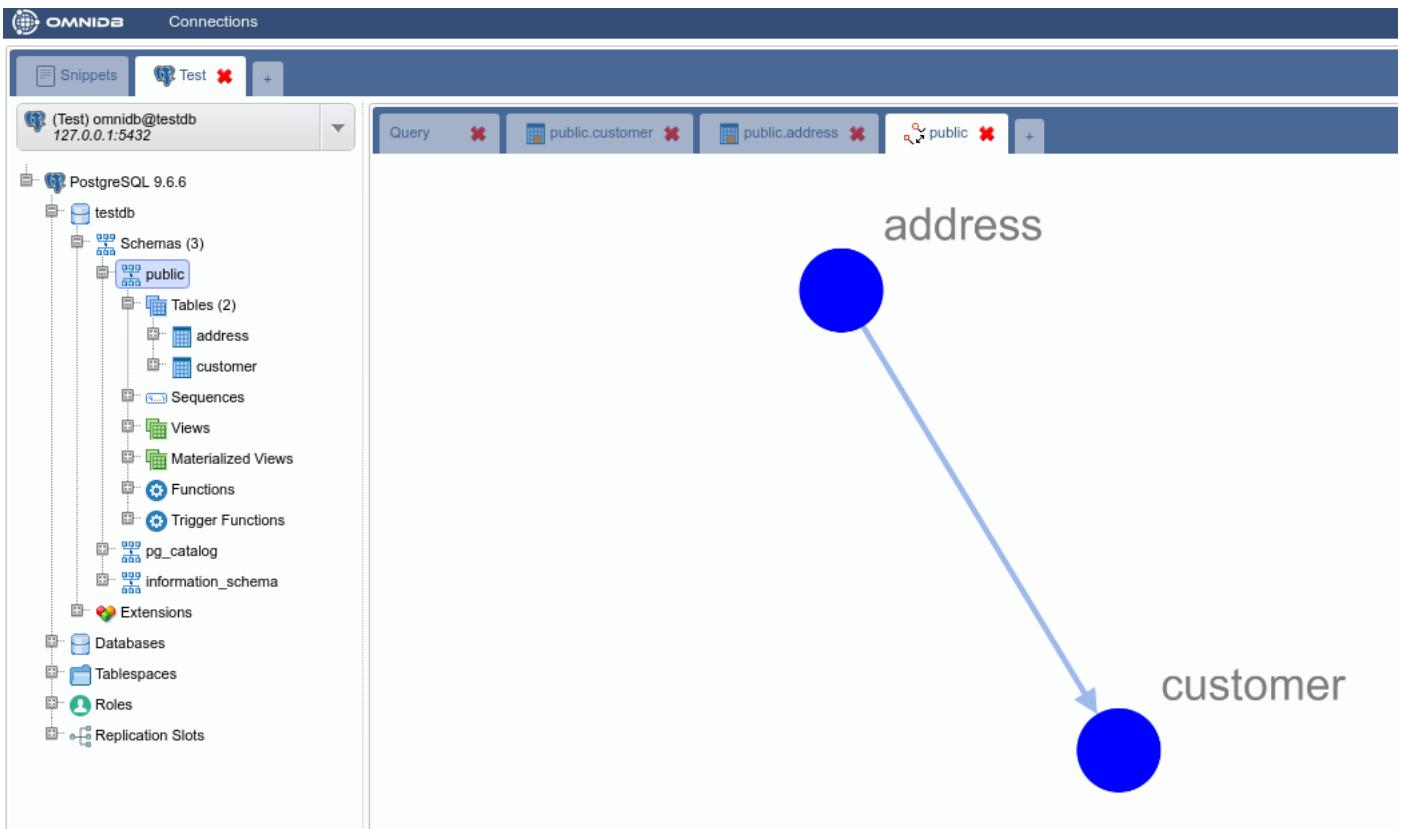
PostgreSQL 9.6.6

- testdb
 - Schemas (3)
 - public
 - Render Graph
 - Simple Graph
 - Complete Graph
 - Alter Schema
 - Drop Schema
 - Sequences
 - Views
 - Materialized Views
 - Functions
 - Trigger Functions
- pg_catalog

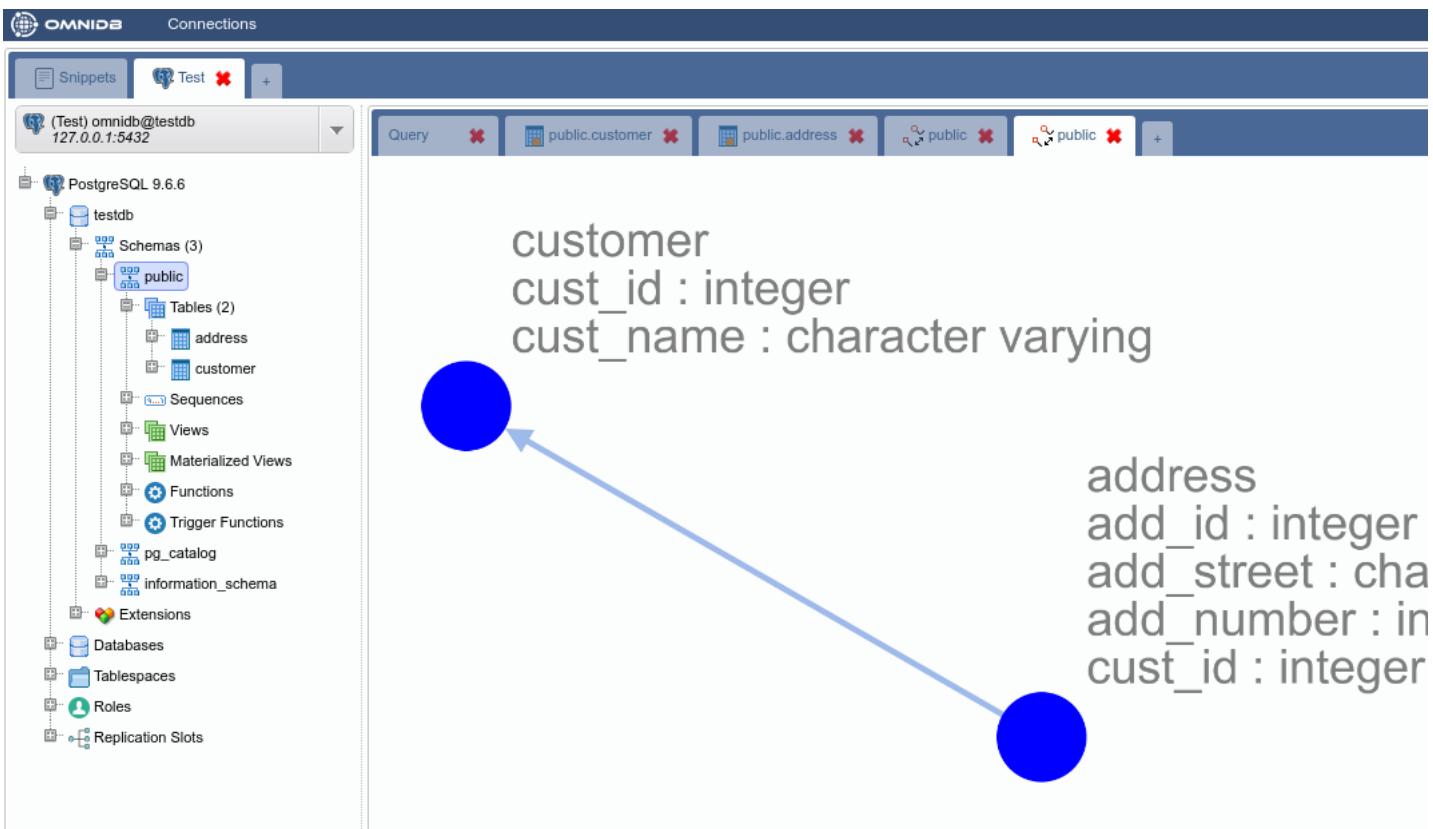
Table Name: address

Columns Constraints Indexes

	Name	Data Typ
1	add_id	integer
2	add_street	character varying(2
3	add_number	integer
4	cust_id	integer
5		



And this is what the *Complete Graph* looks like:



5.0.0.2 Editing tables

OmniDB also lets you edit existing tables (always following DBMS limitations). To test this feature we will add a new column to the table *Customer*. To access the alter table interface just right click the table node and select the action *Table Actions > Alter Table*:

Add the column `cust_age` and save:

	Column Name	Data Type	Nullable	
1	<code>cust_id</code>	integer	NO	✗
2	<code>cust_name</code>	character varying(100)	NO	✗
3	<code>cust_age</code>	integer	YES	✗
4				

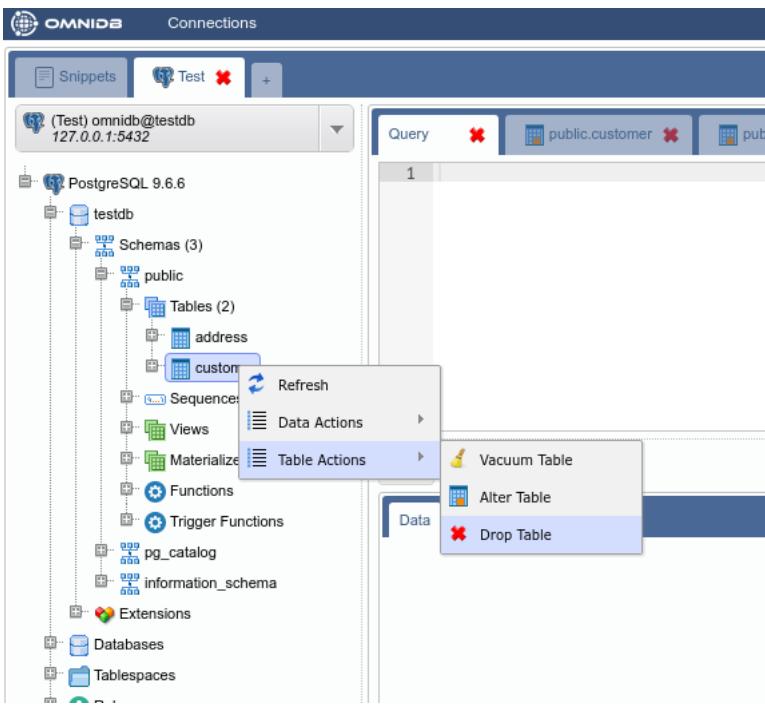
The interface is capable of detecting errors that may occur during alter table operations, showing the command and the error that occurred. To demonstrate it we will try to add the column `cust_name`, which already belongs to this table:

Command: alter table public.customer add column cust_name character varying(100) not null

Message: column "cust_name" of relation "customer" already exists

5.0.0.3 Removing tables

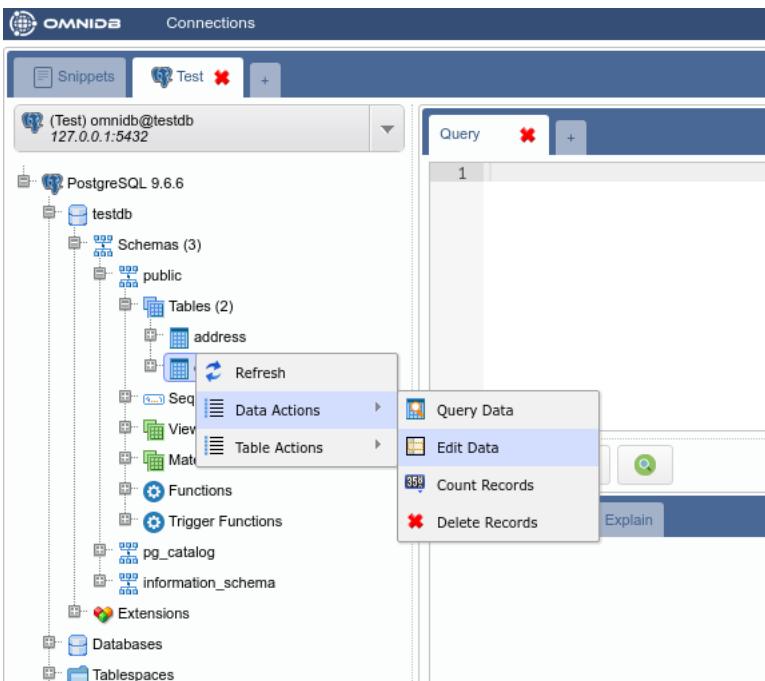
In order to remove a table just right click the table node and select the action `Table Actions > Drop Table`:



6 Managing Table Data

The tool allows us to edit records contained in tables through a very simple and intuitive interface. Given that only a few DBMS have unique identifiers for table records, we opted to allow data editing and removal only for tables that have a primary key. Tables that do not have it can only receive new records.

To access the record editing interface, right click the table node and select the action *Data Actions > Edit Data*:



The screenshot shows the OmniDB interface. On the left, there's a sidebar with a tree view of the database structure. Under 'PostgreSQL 9.6.6', the 'testdb' database is selected, showing 'Schemas (3)', 'Tables (2)' containing 'address' and 'customer', and other objects like 'Sequences', 'Views', etc. The main area has a 'Query' tab open with the SQL command: 'select * from public.customer t'. Below it, the results grid shows one row: '1 order by t.cust_id'. The grid has columns: 'Number of records: 0', 'Response time: 0.017 seconds', and the data itself which includes a primary key icon next to 'cust_id'.

The interface has a SQL editor where you can filter and order records. To prevent that the interface requests too many records, there is a field that limits the number of records to be displayed. The records grid has column names and data types. Columns that belong to the primary key have a key icon next to their names.

The row of the grid that have the symbol * is the row to add new records. Let us insert some records in the table Customer:

This screenshot shows the same OmniDB interface as before, but now with a query to insert records. The SQL command is: 'select * from public.customer t'. Below it, the results grid shows 10 rows of data inserted, each with a primary key icon next to 'cust_id'. The grid has columns: 'Number of records: 0', 'Response time: 0.017 seconds', and the data itself.

	cust_id (integer)	cust_name (character varying)	cust_age (integer)
1	0	Pedro	22
2	1	Ryan	18
3	2	William	23
4	3	Susan	31
5	4	Nicole	19
6	5	Ricardo	45
7	6	Ademar	60
8	7	Felipe	29
9	8	Rafael	21
10	+		

After saving, the records will be inserted and can be edited (only because this table has a primary key). Let's change the *cust_name* of some of the existing records:

This screenshot shows the same OmniDB interface after saving the previous inserts. The SQL command is still 'select * from public.customer t'. The results grid shows the 10 inserted rows. Some rows have been edited: row 3 has 'William' changed to 'William Changed', row 7 has 'Ademar' changed to 'Ademar Changed', and row 8 has 'Felipe' selected (indicated by a blue border).

	cust_id (integer)	cust_name (character varying)	cust_age (integer)
1	0	Pedro	22
2	1	Ryan	18
3	2	William Changed	23
4	3	Susan	31
5	4	Nicole	19
6	5	Ricardo	45
7	6	Ademar Changed	60
8	7	Felipe	29
9	8	Rafael	21
10	+		

Tables can have fields with values represented by very long strings. To help edit these fields, OmniDB has an interface that can be accessed by right clicking the specific cell:

The screenshot shows a database interface with a query window containing the following SQL code:

```
select * from public.customer t
1 order by t.cust_id
```

Below the query window is a table with 10 rows of data. The columns are labeled: cust_id (integer), cust_name (character varying), and cust_age (integer). The data is as follows:

	cust_id (integer)	cust_name (character varying)	cust_age (integer)
1	0	Pedro	22
2	1	Ryan	18
3	2	William Changed	23
4	3	Susan	31
5	4	Nicole	19
6	5	Ricardo	45
7	6	Ademar Changed	60
8	7	Felipe	
9	8	Rafael	21
10	+		

A tooltip "Edit Content" is visible over the Felipe row.

The screenshot shows the OMNIDB application interface. On the left is a tree view of connections and databases. In the center is a large white area for running queries. At the top of this area, there is a command history entry:

```
1 Ademar Changed
```

The interface detects errors that may occur during operations related to records. To demonstrate, let us insert two records with existing cust_id (primary key):

The screenshot shows the OMNIDB interface with the following error messages displayed:

Command: insert into public.customer (cust_id, cust_name, cust_age) values (8, 'Wrong 2', 1)

Message:

```
duplicate key value violates unique constraint "cust_pk"
DETAIL: Key (cust_id)=(8) already exists.
```

Command: insert into public.customer (cust_id, cust_name, cust_age) values (7, 'Wrong 1', 1)

Message:

```
duplicate key value violates unique constraint "cust_pk"
DETAIL: Key (cust_id)=(7) already exists.
```

It shows which commands tried to be executed and the respective errors.

To complete this chapter, let's add some records to the *Address* table:

The screenshot shows the OmniDB interface with three tabs at the top: 'Query', 'public.customer', and 'public.address'. The 'public.address' tab is active. Below it is a query editor with the following code:

```
select * from public.address t
  1 order by t.add_id
```

Below the editor is a table titled 'Query 10 rows' with the following data:

	add_id (integer)	add_street (character varying)	add_number (integer)	cust_id (integer)
1	0	Blue Street	114	0
2	1	Red Street	471	1
3	2	Black Street	355	2
4	3	White Street	1002	3
5	4	Green Street	1056	4
6	5	Purple Street	19	5
7	6	Orange Street	47	6
8	7	Yellow Street	33	7
9	8	Brown Street	29	8
10	+			

7 Writing SQL Queries

The tool comes with a tab system where each tab contains a SQL editor, an action button, an indent button, a field to select the type of command and a space to display the result.

The SQL editor has a feature that helps a lot when creating new queries: SQL code completion. With this feature it is possible to autocomplete columns contained in a table referenced by an alias. To open the autocomplete interface you just have to type the alias and then the dot character:

The screenshot shows the OmniDB interface with a query editor containing the following code:

```
1 select *
2 from customer cust
3 where cust.cust_age > 30
4 order by cust|
```

A tooltip appears over the cursor position, showing the available columns for the 'cust' alias:

- cust.
- cust.cust_id int4
- cust.cust_name varchar
- cust.cust_age int4

Below the editor are buttons for play, refresh, and search, followed by tabs for 'Data', 'Messages', and 'Explain'.

Besides autocompleting table columns the editor also searches for columns contained in subqueries:

The screenshot shows the OmniDB interface with a query editor containing the following code:

```
1 select *
2 from (select cust.cust_name,
3             (select count(*)
4              from address addr
5                 where addr.cust_id = cust.cust_id) as num_addresses
6           from customer cust) subquery
7 where subquery|
```

A tooltip appears over the cursor position, showing the available columns for the 'subquery' alias:

- subquery.
- subquery.cust_name varchar
- subquery.num_addresses int8

Below the editor are buttons for play, refresh, and search, followed by tabs for 'Data', 'Messages', and 'Explain'.

If the query raises an error, OmniDB will show the error message in the *Data* tab and the cursor will be placed in the position indicated by the error message:

```

1 select *
2 from (select cust.cust_name,
3             (select count(*)
4                  from address addr
5                 where addr.cust_id = cust.cust_id) as num_addresses
6          from customer cust) subquery
7 where subquery.cust_name = 'William'

```

Start time: 12/05/2017 08:26:37 Duration: 8.91 ms

Data Messages Explain

column subquery.cust_name does not exist
LINE 7: where subquery.cust_name = 'William'
HINT: Perhaps you meant to reference the column "subquery.cust_name".

When the query is successful, OmniDB shows the number of records returned by the query, the start time and the duration of the query. The *Data* tab will show a data grid with the records records returned by the query.

```

1 select *
2 from (select cust.cust_name,
3             (select count(*)
4                  from address addr
5                 where addr.cust_id = cust.cust_id) as num_addresses
6          from customer cust) subquery

```

Number of records: 9
Start time: 12/05/2017 08:24:52 Duration: 9.519 ms

Data Messages Explain

	cust_name	num_addresses
1	Pedro	1
2	Ryan	1
3	Susan	1
4	Nicole	1
5	Ricardo	1
6	Felipe	1
7	Rafael	1
8	William Changed	1
9	Ademar Changed	1

Just like in the record editing interface each cell can be visualized separately by right clicking it.

If you click on the *Indent SQL* button, OmniDB will reorganize the SQL text to make it prettier:

```

1 select *
2 from
3   (select cust.cust_name,
4           (select count(*)
5              from address addr
6             where addr.cust_id = cust.cust_id) as num_addresses
7      from customer cust) subquery
8 where subquery.cust_name = 'Rafael'

```

Number of records: 1
Start time: 12/05/2017 09:21:01 Duration: 13.114 ms

Data Messages Explain

	cust_name	num_addresses
1	Rafael	1

8 Visualizing Query Plans

OmniDB 2.2.0 introduced a very useful feature: graphical query plan visualization. This may come in handy when writing or optimizing queries, since it allows you to easily identify performance bottlenecks in your SQL query.

For this feature, *SQL Query* inner tab shows 2 buttons: *Explain* (magnifier in orange circle button) and *Explain Analyze* (magnifier in green circle button).

8.0.0.1 Textual visualization

When you click the *Explain* button, OmniDB will execute an EXPLAIN command in your query. Initial visualization is *textual* and will show exactly the output of the EXPLAIN command, but with colored bars representing the estimated cost. The higher the cost, the darker the bar.

The screenshot shows the OmniDB interface with the 'Explain' tab selected. The query window contains the following SQL code:

```
1 select *
2 from
3   (select cust.cust_name,
4    5     (select count(*)
6      from address addr
7      where addr.cust_id = cust.cust_id) as num_addresses
8    from customer cust) subquery
9 where subquery.cust_name = 'Rafael'
```

Below the query window, the status bar shows 'Start time: 12/05/2017 09:52:32 Duration: 9.481 ms'. The toolbar includes icons for Run, Stop, Refresh, and Explain.

The main area displays the 'QUERY PLAN' with numbered steps:

- # Seq Scan on customer cust (cost=0.00..38.27 rows=2 width=226)
- 2 Filter: ((cust_name)::text = 'Rafael'::text)
- 3 SubPlan 1
- 4 Aggregate (cost=12.13..12.14 rows=1 width=8)
- 5 Seq Scan on address addr (cost=0.00..12.12 rows=1 width=0)
- 6 Filter: (cust_id = cust.cust_id)

When you click the *Explain Analyze* button, OmniDB will execute an EXPLAIN ANALYZE command in your query. Beware that this command will really execute the query. Also, the textual visualization will show much more information, and the costs are not estimated as in those provided by the EXPLAIN command; they are real costs.

The screenshot shows the OmniDB interface with the 'Explain' tab selected. The query window contains the same SQL code as the previous screenshot.

Below the query window, the status bar shows 'Start time: 12/05/2017 10:03:26 Duration: 15.073 ms'. The toolbar includes icons for Run, Stop, Refresh, and Explain.

The main area displays the 'QUERY PLAN' with numbered steps and execution details:

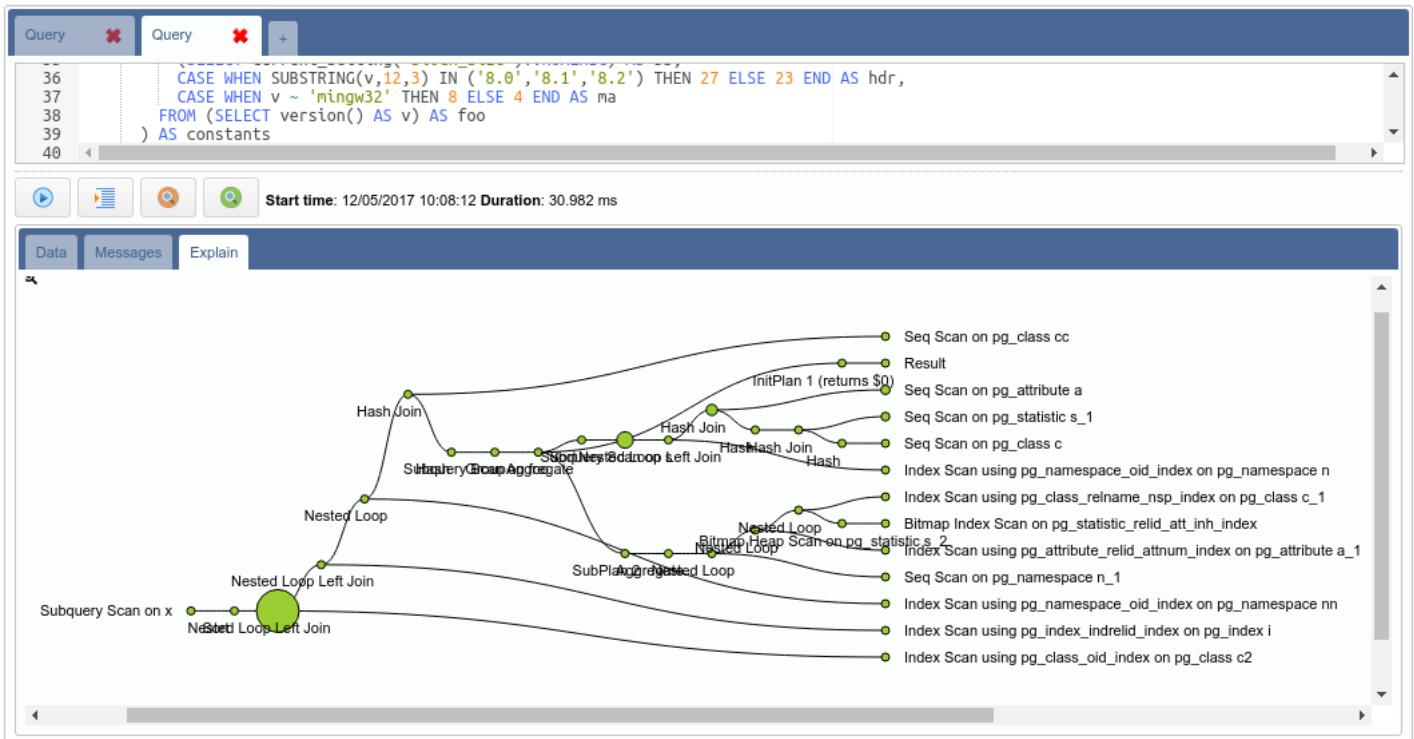
- # Seq Scan on customer cust (cost=0.00..38.27 rows=2 width=226) (actual time=0.029..0.029 rows=1 loops=1)
- 2 Filter: ((cust_name)::text = 'Rafael'::text)
- 3 Rows Removed by Filter: 8
- 4 SubPlan 1
- 5 Aggregate (cost=12.13..12.14 rows=1 width=8) (actual time=0.017..0.017 rows=1 loops=1)
- 6 Seq Scan on address addr (cost=0.00..12.12 rows=1 width=0) (actual time=0.013..0.014 rows=1 loops=1)
- 7 Filter: (cust_id = cust.cust_id)
- 8 Rows Removed by Filter: 8
- 9 Planning time: 0.323 ms
- 10 Execution time: 0.097 ms

8.0.0.2 Tree visualization

Both *Explain* and *Explain Analyze* modes also can graphically represent the textual output into a *tree* diagram. Each circle represent a node executed by the query plan, and the larger the circle, the higher the cost.



When queries become more and more complex, also its query plan can be very complex. With such queries (like the *check bloat* query we executed below) the tree visualization can be very interesting:



The query plan visualization component allows you to easily switch between textual and 2 tree visualizations, which can be zoomed in and out.

9 Visualizing Data

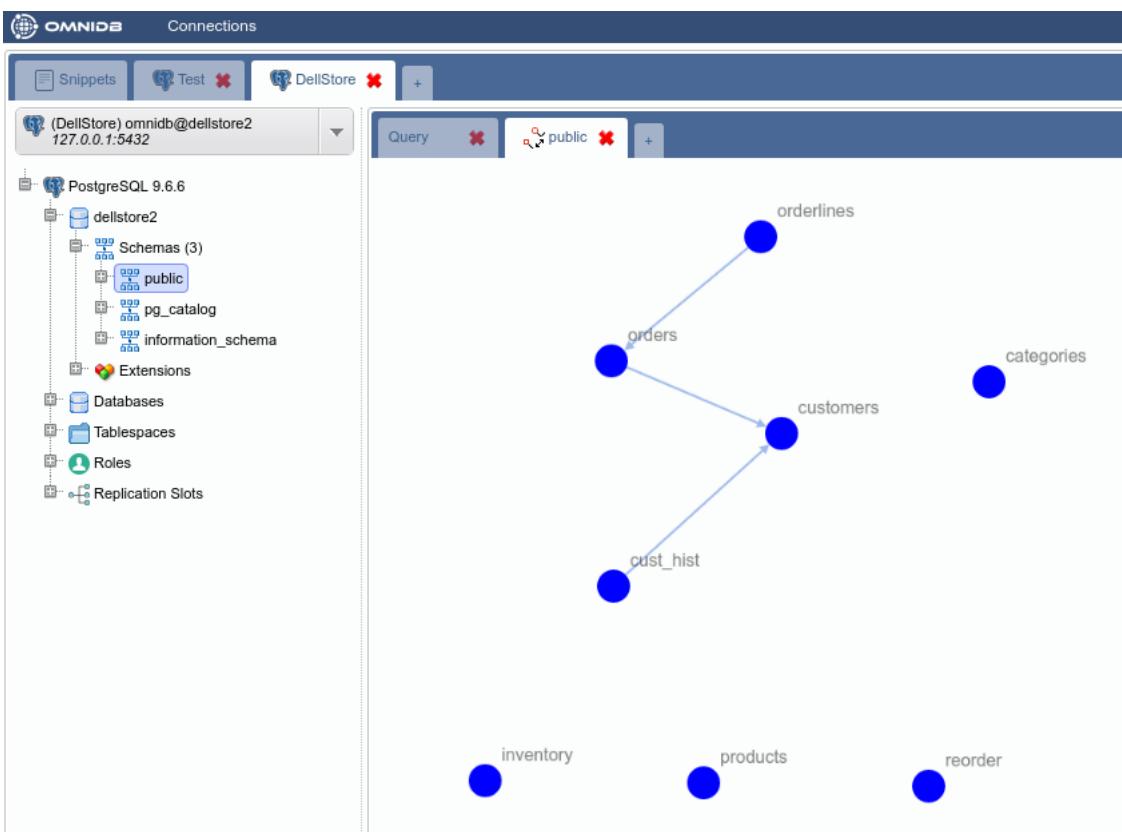
This feature displays a graph with nodes representing tables and edges representing table relationships with foreign keys. Using the mouse, the user is able to zoom in, zoom out, and drag and drop nodes to change its position.

There are two types of graphs: *Simple Graph* and *Complete Graph*.

9.0.0.1 Simple graph

To access it just right click the root node of the tree and then select the action *Render Graph > Simple Graph*:

The screenshot shows the OMNIDB interface with a PostgreSQL 9.6.6 connection named 'dellstore2'. In the left sidebar, under the 'Schemas' section for 'dellstore2', there is a context menu with three options: 'Render Graph', 'Simple Graph', and 'Complete Graph'. The 'Simple Graph' option is highlighted with a light blue background. The main query editor window is empty, showing a single row labeled '1'.



9.0.0.2 Complete graph

This graph displays tables with all its columns and respective data types. Additionally, edges now are labeled with information about the specific foreign key. To access it just right click the root node of the tree and then select the action *Render Graph > Complete Graph*:

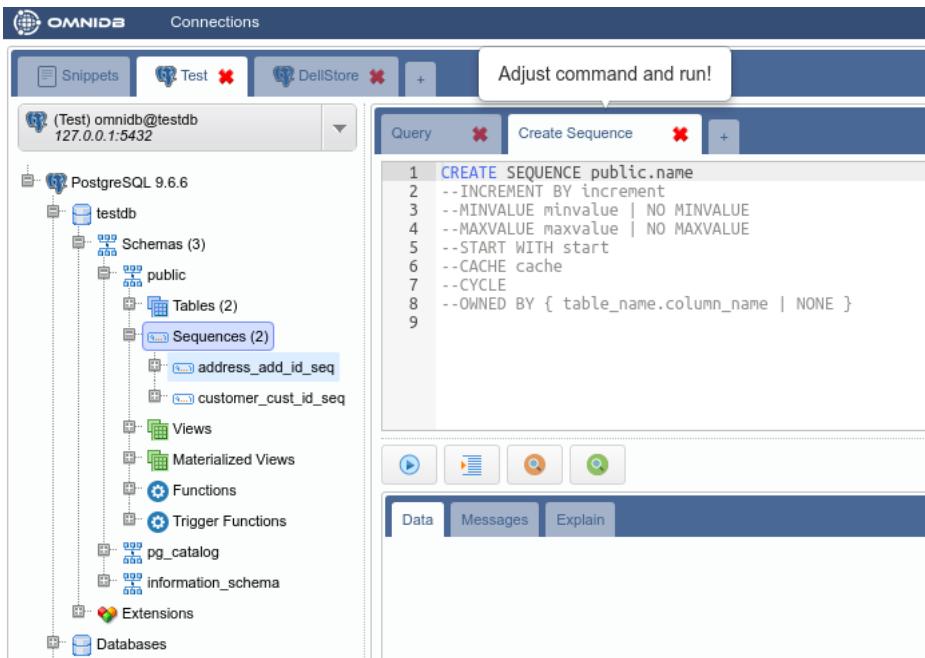
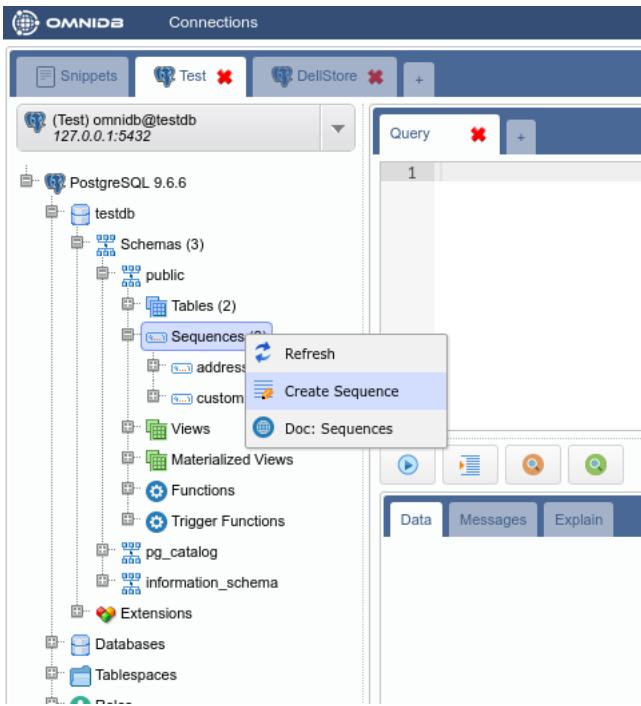
The screenshot shows the OMNIDB interface with a connection to 'DellStore' on port 5432. In the left sidebar, under 'dellstore2', there is a 'Schemas (3)' section containing 'public'. A context menu is open over the 'public' schema, with 'Render Graph' selected. The graph view on the right shows a single node labeled 'order' with three outgoing edges.

The screenshot shows the OMNIDB interface with a connection to 'DellStore' on port 5432. In the left sidebar, under 'dellstore2', there is a 'Schemas (3)' section containing 'public', 'pg_catalog', and 'information_schema'. A context menu is open over the 'public' schema, with 'Simple Graph' selected. The graph view on the right displays nodes for various tables: 'categories', 'customers', 'orders', 'products', 'reorder', 'sales', and 'inventory'. Each node contains a detailed list of its columns.

10 Managing other PostgreSQL Elements

With the exception of tables, all PostgreSQL structures are possible to be managed with the use of *SQL templates*. This gives the user more power than using graphical forms to manipulate structures.

For example, let's consider the sequences inside the schema `public` of the database `testdb`. To create a new sequence, right click on the `Sequences` node, and choose *Create Sequence*.



After you change the name of the sequence, you can uncomment other command options and set them accordingly to your needs. When the entire command looks fine, just execute it and the new sequence will be created:

The screenshot shows the OMNIDB PostgreSQL interface. On the left, the database structure for 'testdb' is displayed, including 'Schemas (3)', 'Tables (2)', 'Sequences (3)', 'Views', 'Materialized Views', 'Functions', and 'Trigger Functions'. A right-click context menu is open over the 'seq_test' sequence, showing options like 'Alter Sequence', 'Drop Sequence', and 'Copy'. The main panel on the right shows a query editor with the following SQL code:

```
1 CREATE SEQUENCE public.seq_test
2 INCREMENT BY 2
3 -MINVALUE minvalue | NO MINVALUE
4 -MAXVALUE maxvalue | NO MAXVALUE
5 -START WITH start
6 -CACHE cache
7 -CYCLE
8 -OWNED BY { table_name.column_name | NONE }
```

Below the query editor, the status bar indicates 'Start time: 12/05/2017 11:37:06 Duration: 116.53 ms'. At the bottom, tabs for 'Data', 'Messages', and 'Explain' are visible, along with a 'Done.' message.

With right click on an existing sequence, you can alter or drop it. It will work the same way as the creation, by using a SQL template for the user to change.

This screenshot shows the same PostgreSQL interface as above, but with a different sequence selected ('seq_1'). A right-click context menu is open over 'seq_1', showing options like 'Alter Sequence', 'Drop Sequence', and 'Copy'. The 'Alter Sequence' option is highlighted. The database structure on the left remains the same, showing 'testdb' with its various schema objects.

```

1 ALTER SEQUENCE public.seq_test
2 INCREMENT BY 1
3 --MINVALUE minvalue | NO MINVALUE
4 --MAXVALUE maxvalue | NO MAXVALUE
5 --START WITH start
6 --RESTART
7 --RESTART WITH restart
8 --CACHE cache
9 --CYCLE
10 --NO CYCLE
11 --OWNED BY { table_name.column_name | NONE }
12 --OWNER TO { new_owner | CURRENT_USER | SESSION_USER }
13 --RENAME TO new_name
14 --SET SCHEMA new_schema
15

```

Start time: 12/05/2017 11:39:39 Duration: 106.117 ms

Data Messages Explain

Done.

```

1 DROP SEQUENCE public.seq_test
2 --CASCADE
3

```

Start time: 12/05/2017 11:40:08 Duration: 115.908 ms

Data Messages Explain

Done.

11 Additional Features

11.0.0.1 SQL History

Every interaction the user does with every database is logged in OmniDB's *SQL History*. To access it you need to click on the icon with an *H* on the upper right corner. OmniDB will open an inner tab with all actions in a paginated grid.

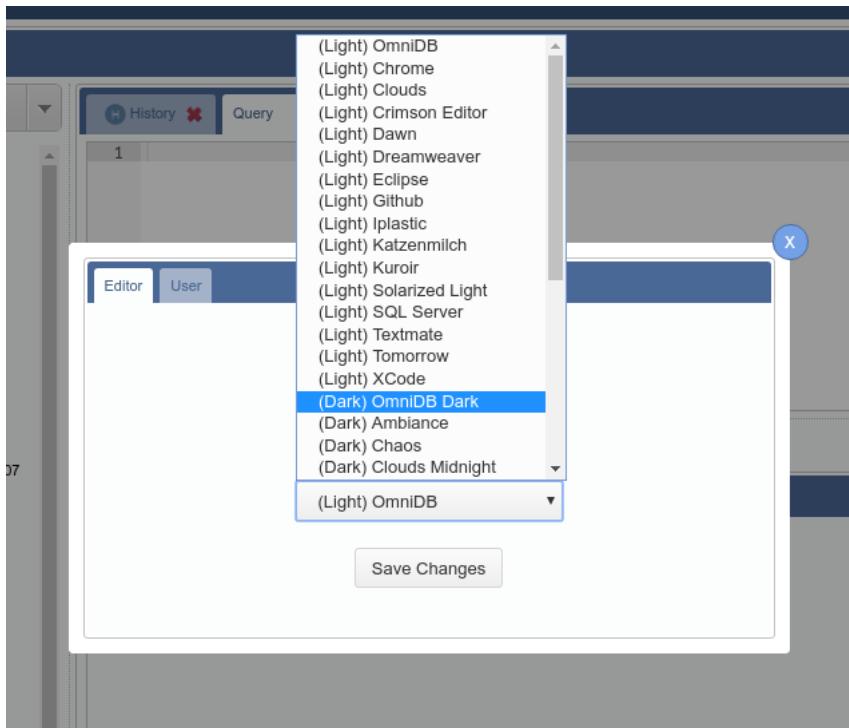
	Start	End	Duration	Status	C
1	2017-12-05 13:40:08.201896	2017-12-05 13:40:08.317804	115.908 ms	Green	DROP SEQUENCE public.seq_test .
2	2017-12-05 13:39:39.692279	2017-12-05 13:39:39.798396	106.117 ms	Green	ALTER SEQUENCE public.seq_test
3	2017-12-05 13:37:06.486254	2017-12-05 13:37:06.602784	116.53 ms	Green	CREATE SEQUENCE public.seq_te
4	2017-12-05 11:48:17.332837	2017-12-05 11:48:17.430137	97.3 ms	Green	explain select * from (select cust.cus
5	2017-12-05 11:21:01.471111	2017-12-05 11:21:01.484225	13.114 ms	Green	select * from (select cust.cust_name,
6	2017-12-05 11:20:45.249422	2017-12-05 11:20:45.258836	9.414 ms	Green	select * from (select cust.cust_name,
7	2017-12-05 11:20:33.822952	2017-12-05 11:20:33.916459	93.507 ms	Green	-- Querying Data select t.* from publi
8	2017-12-05 11:20:26.960773	2017-12-05 11:20:26.971530	10.757 ms	Green	select * from (select cust.cust_name,
9	2017-12-05 11:13:30.014418	2017-12-05 11:13:30.024597	10.179 ms	Red	select * from (select cust.cust_name,
10	2017-12-05 10:26:37.547517	2017-12-05 10:26:37.556427	8.91 ms	Red	select * from (select cust.cust_name,
11	2017-12-05 10:24:52.766592	2017-12-05 10:24:52.776111	9.519 ms	Green	select * from (select cust.cust_name,
12	2017-12-05 10:22:35.513501	2017-12-05 10:22:35.604137	90.636 ms	Green	select * from (select cust.cust_name,
13	2017-11-30 19:25:50.454071	2017-11-30 19:25:50.472099	18.028 ms	Green	DROP TABLE public.customer --CAS
14	2017-11-30 19:25:46.188620	2017-11-30 19:25:46.217045	28.425 ms	Green	DROP TABLE public.address --CAS
15	2017-11-30 18:40:27.042825	2017-11-30 18:40:28.937599	00:00:01	Red	select pg_sleep(30)
16	2017-11-30 18:39:46.856919	2017-11-30 18:40:16.874995	00:00:30	Green	select pg_sleep(30)
17	2017-11-30 18:39:15.600141	2017-11-30 18:39:20.617678	00:00:05	Green	select pg_sleep(5)
18	2017-11-30 18:39:07.229556	2017-11-30 18:39:12.255480	00:00:05	Green	select pg_sleep(5)
19	2017-11-30 18:30:11.140417	2017-11-30 18:30:11.161083	20.666 ms	Green	-- Querying Data select t.* from publi

Each action shows date time it started, the time it ended, the duration, the status and the command. As every grid in OmniDB, you can right click on the command and click *View Content*, where another pop-up will open showing the content in a larger text editor.

11.0.0.2 User Settings

Also in the upper right corner, by clicking in the gear-like icon, OmniDB will open the *User Settings* pop-up. It is composed by two tabs:

- **User:** Allows the user to change its password. More user settings will be added in the future.
- **Editor:** Allows the user to change the font size of the SQL Editor, and also change the entire OmniDB theme. There are a lot of OmniDB themes, each of them change the syntax highlight color of the editor. They are also categorized in light and dark themes. A light theme is the default; a dark theme will change the entire interface of OmniDB.



Every change in user settings require that you either:

- Refresh the page, if you are using OmniDB server and the interface through a web browser; or
- Open and close OmniDB, if you are using OmniDB app.

A screenshot of the OmniDB interface. On the left is a tree view of a PostgreSQL 9.6.6 database named 'testdb'. Nodes include Schemas (3), public, Tables (2), add, cust, Sequences, Views, Materialized Views, Functions, Trigger Functions, pg_catalog, information_schema, Extensions, Databases, Tablespaces, Roles, and Replication Slots. Right-clicking on the 'add' table node opens a context menu with options: Data Actions (selected), Query Data, Edit Data, Count Records, and Delete Records. The 'Query Data' option is highlighted. The main area shows a 'Query' tab with the SQL command: '1 -- Querying Data 2 select t.* 3 from public.address t'. Below it is a table titled 'Number of records: 9 Start time: 12/05/2017 11:56:09 Duration: 92.266 ms'. The table has columns add_number and cust_id. The data is as follows:

	add_number	cust_id		
1	114	0		
2	1	Red Street	471	1
3	2	Black Street	355	2
4	3	White Street	1002	3
5	4	Green Street	1056	4
6	5	Purple Street	19	5
7	6	Orange Street	47	6
8	7	Yellow Street	33	7
9	8	Brown Street	29	8

11.0.0.3 Contextual Help

Most of tree nodes (generally grouping ones like *Schemas* or *Tables*) offer contextual help. This feature can be accessed by right-clicking the tree node. When you click in the *Doc: ...* option, OmniDB will open an inner tab showing a web browser pointing to the specific page in the online **PostgreSQL Documentation**. Also, it will redirect to the specific page considering the PostgreSQL version you are connected to.

The screenshot shows the OmniDB interface with a PostgreSQL connection named 'testdb' selected. The left sidebar shows the database structure with 'testdb' expanded, displaying 'Schemas (3)' which includes 'public', 'pg_catalog', and 'information_schema'. A context menu is open over the 'Schemas' entry, with 'Create Schema' highlighted. To the right, a browser window displays the PostgreSQL documentation for 'PostgreSQL 9.6.6 Documentation' under 'Chapter 5. Data Definition'. The page title is '5.8. Schemas'. A note in the documentation states: 'Note: Users of a cluster do not necessarily have the privilege to access every database in the cluster. Sharing of user names means that there cannot be different users named, say, joe in two databases in the same cluster; but the system can be configured to allow joe access to only some of the databases.' Navigation links include 'Home', 'Documentation', 'Manuals', 'PostgreSQL 9.6.6', 'Prev', and 'Up'.

12 OmniDB Config Tool

Every installation of OmniDB also comes with a small CLI utility called *OmniDB Config*. It will have a different file name, depending on the way you installed OmniDB:

- If you are using a tarball or zip package, it is called **omnidb-config**, for both server and app versions;
- If you used an installer (like the .deb file) of server version, it is called **omnidb-config-server**;
- If you used an installer of app version, it is called **omnidb-config-app**.

Despite having different names, the utility does exactly the same. If you used an installer, it will be put in your \$PATH.

```
user@machine:~$ omnidb-config-app --help
Usage: omnidb-config-app [options]

Options:
  --version           show program's version number and exit
  -h, --help          show this help message and exit
  -c username password, --createsuperuser=username password
                      create super user: -c username password
  -a, --vacuum        databases maintenance
  -r, --resetdatabase reset databases
```

12.0.0.1 Create super user

Option -c allows you to create a new super user, without needing to open OmniDB interface.

```
user@machine:~$ omnidb-config-app -c william password
Creating superuser...
Superuser created.
```

12.0.0.2 Vacuum

OmniDB has two databases:

- **omnidb.db**: Stores all users and connections, and other OmniDB related stuff;
- **Sessions database**: Stores Django user sessions.

Both databases are SQLite, so it can be useful to vacuum them sometimes to reduce file size. This can be done with the -a option.

```
user@machine:~$ omnidb-config-app -a
Vacuuming OmniDB database...
Done.
Vacuuming Sessions database...
Done.
```

12.0.0.3 Reset database

If you wish to wipe out all OmniDB information and get a clean database as it was just installed, you can use the `-r` option. Use it with caution!

```
user@machine:~$ omnadb-config-app -r
*** ATENTION *** ALL USERS DATA WILL BE LOST
Would you like to continue? (y/n) y
Cleaning users...
Done.
Cleaning sessions...
Vacuuming OmniDB database...
Done.
Vacuuming Sessions database...
Done.
```

13 Writing and Debugging PL/pgSQL Functions

13.0.0.1 Introduction

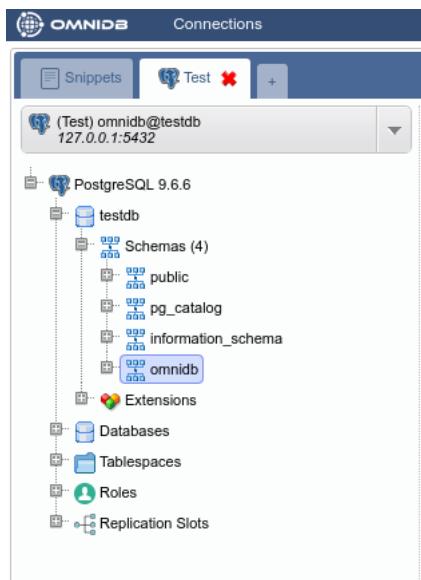
PostgreSQL is more than a RDBMS engine. It is a developing platform. It provides a very powerful and flexible programming language called PL/pgSQL. Using this language you can write your own *user-defined functions* to achieve abstraction levels and procedural calculations that would be difficult to achieve with plain SQL (and sometimes impossible to achieve without context-switching with the application). While you always could develop and manage your own functions within OmniDB, it is a recent feature that allows you to also *debug* your own functions.

OmniDB 2.3.0 introduces this great feature: a debugger for PL/pgSQL functions. It was implemented by scratch and takes advantage of hooks, an extensibility in PostgreSQL's source code that allows us to perform custom actions when specific events are triggered in the database. For the debugger we use hooks that are triggered when PL/pgSQL functions are called, and each statement is executed.

This requires the user to install a binary library called `omnidb_plugin` and enable it in PostgreSQL's config file. The debugger also uses a special schema with special tables to control the whole debugging process. This can be manually created or with an extension.

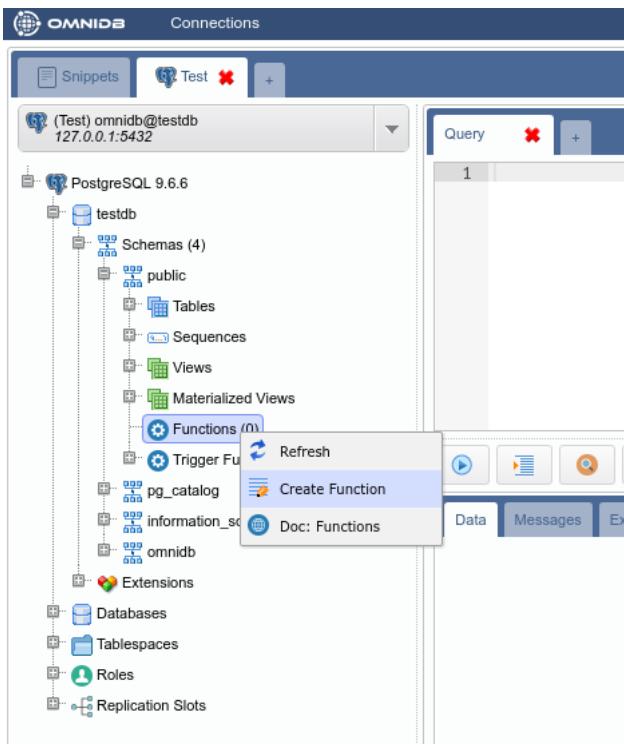
For more details on the installation, please refer to the [instructions](#). Also please read the notes in this document, to be aware that currently there are some limitations.

After successfully installing the debugger, you will see a schema called `omnidb` in your database. Also, if you compiled the debugger yourself, you can install it as a PostgreSQL extension, and in this case it will appear under the *Extensions* tree node.



13.0.0.2 Writing functions

In the `public` schema, right-click the `Functions` node and click on *Create Function*. It will open a *SQL Query* inner tab, already containing a SQL Template to help you create your first PL/pgSQL function.



Adjust command and run!

```

1 CREATE OR REPLACE FUNCTION public.name
2   --( [ argmode ] [ argname ] argtype [ { DEFAULT | = } default_expr ]
3   --)
4   --RETURNS retype
5   --RETURNS TABLE ( column_name column_type )
6 LANGUAGE plpgsql
7   --IMMUTABLE | STABLE | VOLATILE
8   --STRICT
9   --SECURITY DEFINER
10  --COST execution_cost
11  --ROWS result_rows
12 AS
13 $function$

```

Function public.name created.

You can refer to PostgreSQL documentation on how to write user-defined functions. No need to open a new browser tab: just right-click the *Functions* node and click on *Doc: Functions* to view the documentation inside OmniDB.

For now, let us replace this SQL template entirely for the source code below:

```

CREATE OR REPLACE FUNCTION public.fnc_count_vowels (p_input text)
RETURNS integer LANGUAGE plpgsql AS
$function$
DECLARE
    str text;
    ret integer;
    i integer;
    len integer;
    tmp text;
BEGIN
    str := upper(p_input);
    ret := 0;
    i := 1;
    len := length(p_input);
    WHILE i <= len LOOP
        IF substr(str, i, 1) IN ('A', 'E', 'I', 'O', 'U') THEN
            SELECT pg_sleep(1) INTO tmp;
            ret := ret + 1;
        END IF;
        i := i + 1;
    END LOOP;

```

```

RETURN ret;
END;
$function$

```

This will create a function called `fnc_count_vowels` inside the schema `public`. This function takes a text argument called `p_input` and counts how many vowels there are in this `string`. Then returns this count.

To create the function, execute the command in the SQL Query inner tab. If successful, the function will appear under the *Functions* tree node (you can refresh it by right-clicking and then clicking in *Refresh*). By expanding the function node as well, you can see its return type and its argument.

```

1 CREATE OR REPLACE FUNCTION public.fnc_count_vowels (p_input text)
2 RETURNS integer LANGUAGE plpgsql AS
3 $function$
4 DECLARE
5   str text;
6   ret integer;
7   i integer;
8   len integer;
9   tmp text;
10 BEGIN
11   str := upper(p_input);
12   ret := 0;
13   i := 1;
14   len := length(p_input);
15   WHILE i <= len LOOP
16     IF substr(str, i, 1) in ('A', 'E', 'I', 'O', 'U') THEN
17       SELECT pg_sleep(1) INTO tmp;
18       ret := ret + 1;
19     END IF;
20     i := i + 1;
21   END LOOP;
22   RETURN ret;
23 END;
24 $function$

```

Now let us execute this new function for the first time. Open a simple SQL Query inner tab and execute the following SQL query:

```
SELECT public.fnc_count_vowels('The quick brown fox jumps over the lazy dog.')
```

	fnc_count_vowels
1	11

Note how the query returns a single value, containing the number of vowels in the text. Note also how the query took several seconds to finish; this is caused by the `pg_sleep` we put in the source code of the function `fnc_count_vowels`.

By right-clicking the function node, you can see there are actions to edit and drop it. As you probably guessed, each action will open SQL Query inner tabs with handy SQL templates in them. But the most interesting action right now is *Debug Function*. Go ahead and click it!

The screenshot shows the OMNIDB interface for PostgreSQL 9.6.6. On the left, the database tree shows the 'testdb' schema with its various objects like tables, sequences, views, and functions. The 'Functions' section contains one function named 'fnc_count_vowels'. A context menu is open over this function, with the 'Debug Function' option highlighted. The main query editor on the right displays the SQL code for this function:

```

1 SELECT public.fnc_count_vowels

```

13.0.0.3 Debugging functions

The debugger is a specific inner tab composed of a SQL editor that will show the process step by step on top of the function source code, and 5 tabs to manage and view different parts of the debugger.

The screenshot shows the OMNIDB debugger interface for the 'fnc_count_vowels' function. The SQL editor at the top shows the function's source code:

```

1 DECLARE
2     str text;
3     ret integer;
4     i integer;
5     len integer;
6     tmp text;
7 BEGIN
8     str := upper(p_input);
9     ret := 0;
10    i := 1;
11    len := length(p_input);
12    WHILE i <= len LOOP
13        IF substr(str, i, 1) in ('A', 'E', 'I', 'O', 'U') THEN

```

Below the SQL editor, there is a button labeled "Adjust parameters and start". The bottom part of the interface features five tabs: Parameters, Variables, Result, Messages, and Statistics. The "Parameters" tab is currently active, displaying a table with one row:

	Parameter	Value
1	p_input	text

- Parameters:** Before the debugging process starts, the user must provide all the parameters in this tab. Parameters must be provided exactly the same way you would provide them if you were executing the function in plain SQL, quoting strings for instance;
- Variables:** This grid displays the current value of each variable that exists in the current execution context, it will be updated with every step;
- Result:** When the function ends, this tab will show the result of the function call. It could be empty, a single value or even a set of rows;
- Messages:** Messages returned explicitly by RAISE commands or even automatic messages from PostgreSQL will be presented in this tab;
- Statistics:** At the end of the debugging process, a chart depicting execution times for each line in the function body will be presented in this tab. Additionally, the SQL editor will be updated with a set of colors representing a heat map, from blue to red, according to the max duration of each line.

Now let us start debugging this function. First thing to do is to fill *every* parameter in the *Parameters* tab:

Query Create Function Query Debugger: fnc_count_vowels

```

1 DECLARE
2   str text;
3   ret integer;
4   i integer;
5   len integer;
6   tmp text;
7 BEGIN
8   str := upper(p_input);
9   ret := 0;
10  i := 1;
11  len := length(p_input);
12  WHILE i <= len LOOP
13    IF substr(str, i, 1) IN ('A', 'E', 'I', 'O', 'U') THEN
14      SELECT pg_sleep(1) INTO tmp;
15      ret := ret + 1;
16    END IF;
17  END LOOP;
18  RETURN ret;
19 END;

```

Adjust parameters and start

Parameters Variables Result Messages Statistics

	Parameter	Value
1	p_input	text 'The quick brown fox jumps over the lazy dog.'

Then click on the *Start* button. Note how OmniDB automatically goes to the *Variables* tab, which is the interesting tab now that the function is being debugged. The argument `p_input` is now called `$1`, indicating the first argument of the function. Also note the variable `found`, which is a PostgreSQL reserved variable that indicates whether or not a query has returned values inside of the function.

Also note that OmniDB points to the first line of the source code of the function, highlighting it in green. This is the line that is about to be executed.

Query Create Function Query Debugger: fnc_count_vowels

```

3   str text;
4   ret integer;
5   i integer;
6   len integer;
7   tmp text;
8 BEGIN
9   str := upper(p_input);
10  ret := 0;
11  i := 1;
12  len := length(p_input);
13  WHILE i <= len LOOP
14    IF substr(str, i, 1) IN ('A', 'E', 'I', 'O', 'U') THEN
15      SELECT pg_sleep(1) INTO tmp;
16      ret := ret + 1;
17    END IF;
18  END LOOP;
19  RETURN ret;
20 END;

```

Ready

Parameters Variables Result Messages Statistics

	Variable	Attribute	Type	Value
1	\$1		text	The quick brown fox jumps over the lazy dog.
2	found		bool	f
3	str		text	NULL
4	ret		int4	NULL
5	i		int4	NULL
6	len		int4	NULL
7	tmp		text	NULL

Now click in the first button below the SQL editor. It is the *Step Over* button, and it means that OmniDB will execute the next statement and stop right after it.

Query Create Function Query Debugger: fnc_count_vowels

```

4  ret integer;
5  i  integer;
6  len integer;
7  tmp text;
8 BEGIN
9  str := upper(p_input);
10 ret := 0;
11 i := 1;
12 len := length(p_input);
13 WHILE i <= len LOOP
14   IF substr(str, i, 1) in ('A', 'E', 'I', 'O', 'U') THEN
15     SELECT pg_sleep(1) INTO tmp;
16     ret := ret + 1;
17   END IF;

```

	Variable	Attribute	Type	Value
1	\$1		text	The quick brown fox jumps over the lazy dog.
2	found		bool	f
3	str		text	THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
4	ret		int4	NULL
5	i		int4	NULL
6	len		int4	NULL
7	tmp		text	NULL

Note how the variable `str` has the value assigned to it during execution of line 9. Right now OmniDB is about to execute line 10, showing the current execution state.

Now that you know how to step over, let us speed up things a little bit. Click on the header of the line 20, the last line of code. By doing this, you just placed a *breakpoint*. The debugger interface allows you to place one breakpoint at a time.

Query Create Function Query Debugger: fnc_count_vowels

```

8 BEGIN
9  str := upper(p_input);
10 ret := 0;
11 i := 1;
12 len := length(p_input);
13 WHILE i <= len LOOP
14   IF substr(str, i, 1) in ('A', 'E', 'I', 'O', 'U') THEN
15     SELECT pg_sleep(1) INTO tmp;
16     ret := ret + 1;
17   END IF;
18   i := i + 1;
19 END LOOP;
20 RETURN ret;
21 END;

```

	Variable	Attribute	Type	Value
1	\$1		text	The quick brown fox jumps over the lazy dog.
2	found		bool	f
3	str		text	THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
4	ret		int4	NULL
5	i		int4	NULL
6	len		int4	NULL
7	tmp		text	NULL

After setting a breakpoint, you can click in the second button, *Resume*. OmniDB will carry on with the debugging process until it reaches the line of code with the breakpoint. This may take a while because of the `pg_sleep` commands we put in the source code. Note that if you click this button without previously setting a breakpoint, OmniDB will execute the entire function to the end.

Query Create Function Query Debugger: fnc_count_vowels

```

9   str := upper(p_input);
10  ret := 0;
11  i := 1;
12  len := length(p_input);
13  WHILE i <= len LOOP
14    IF substr(str, i, 1) in ('A', 'E', 'I', 'O', 'U') THEN
15      SELECT pg_sleep(1) INTO tmp;
16      ret := ret + 1;
17    END IF;
18    i := i + 1;
19  END LOOP;
20  RETURN ret;
21 END;
22

```

Cancel **Ready**

Parameters Variables Result Messages Statistics

	Variable	Attribute	Type	Value
1	\$1		text	The quick brown fox jumps over the lazy dog.
2	found		bool	t
3	str		text	THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
4	ret		int4	11
5	i		int4	45
6	len		int4	44
7	tmp		text	

Observe the values for each variable. We can see that the value of `ret` is 11 even before the function finishes. Also note that OmniDB does not remove the breakpoint you placed. To do that, you can click in the breakpoint little icon. Now hit *Resume* again. Let us see now what happens when the function finishes.

Query Create Function Query Debugger: fnc_count_vowels

```

9   str := upper(p_input);
10  ret := 0;
11  i := 1;
12  len := length(p_input);
13  WHILE i <= len LOOP
14    IF substr(str, i, 1) in ('A', 'E', 'I', 'O', 'U') THEN
15      SELECT pg_sleep(1) INTO tmp;
16      ret := ret + 1;
17    END IF;
18    i := i + 1;
19  END LOOP;
20  RETURN ret;
21 END;
22

```

Finished - Total duration: 12.471 s

Parameters Variables Result Messages Statistics

Line Number	Duration (s)
9	0.0144
10	0.0133
11	0.0134
12	0.0141
13	0
14	0.0134
15	1.015
16	0.0135

OmniDB will go automatically to the *Statistics* tab, which shows 2 interesting features:

- *Sum of Duration per Line of Code Chart*: in the bottom, a chart represents total duration of the function distributed in the lines of code. With this chart, you can easily spot bottlenecks in your code. In our example, it was line 15, which we deliberately put a `pg_sleep(1)` call;
- *Colored lines of source code*: OmniDB colors the lines accordingly to the numbers seen in the chart. Colors vary from blue (small duration), passing through yellow (medium duration) until red (high duration), as in a *temperature* diagram.

Also note the *Total duration* message, which shows execution time of the function, without considering the time you spent analyzing it.

13.0.0.4 Inspecting record attribute values

An interesting feature that we do not usually see in other debuggers is the ability to inspect each attribute of a variable of type `record`. OmniDB debugger does that as it is split into different variables, allowing you to see the value and type of each attribute.

To illustrate that, let us create another function, similar to the previous one, but now called `fnc_count_vowels2`:

```

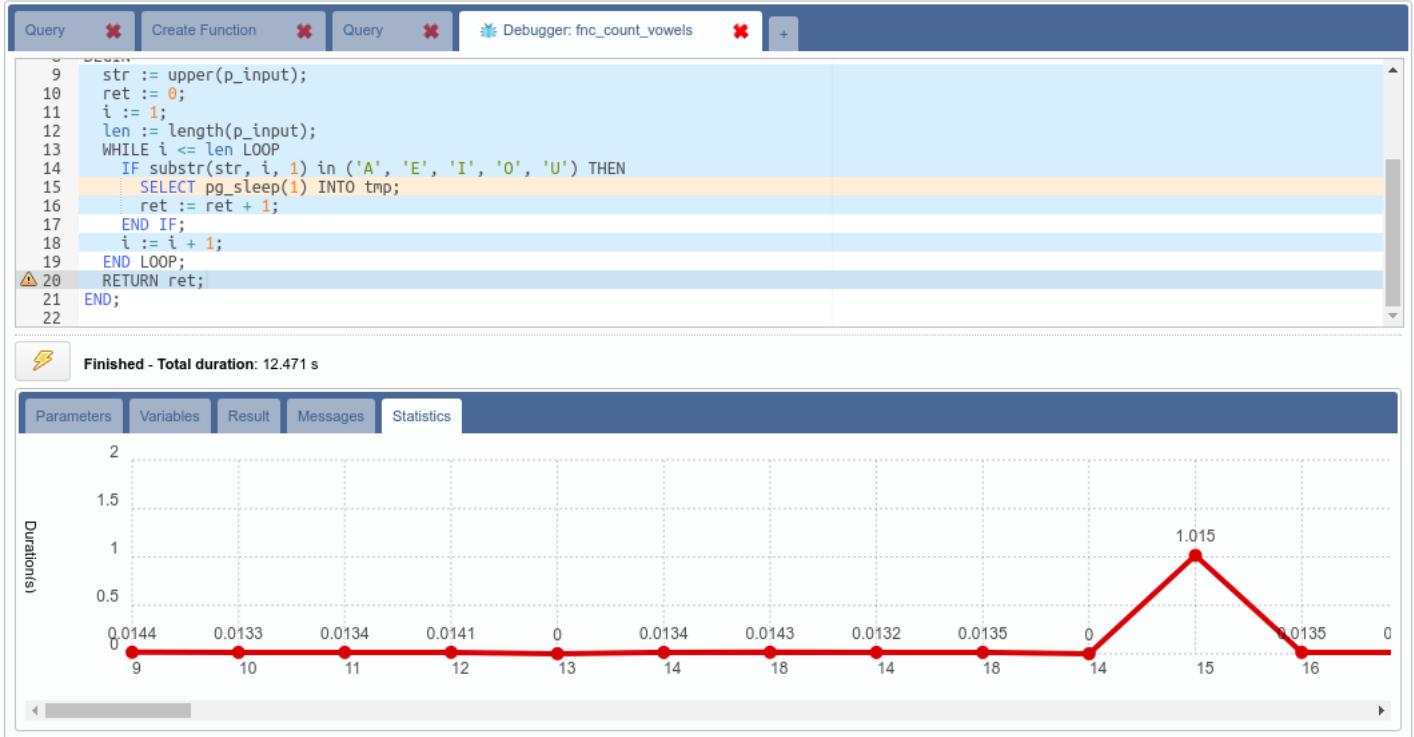
CREATE OR REPLACE FUNCTION public.fnc_count_vowels2 (p_input text)
RETURNS integer LANGUAGE plpgsql AS
$function$
DECLARE
  str text;

```

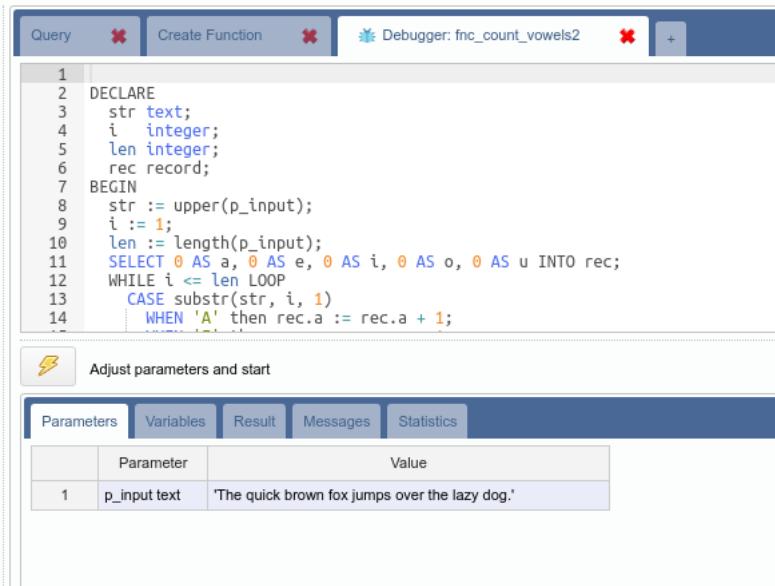
```

i integer;
len integer;
rec record;
BEGIN
str := upper(p_input);
i := 1;
len := length(p_input);
SELECT 0 AS a, 0 AS e, 0 AS i, 0 AS o, 0 AS u INTO rec;
WHILE i <= len LOOP
CASE substr(str, i, 1)
WHEN 'A' then rec.a := rec.a + 1;
WHEN 'E' then rec.e := rec.e + 1;
WHEN 'I' then rec.i := rec.i + 1;
WHEN 'O' then rec.o := rec.o + 1;
WHEN 'U' then rec.u := rec.u + 1;
ELSE NULL;
END CASE;
i := i + 1;
END LOOP;
RETURN rec.a + rec.e + rec.i + rec.o + rec.u;
END;
$function$

```



Observe how we keep track of every vowel count individually. Now let us start debugging it, using the same text as before ('The quick brown fox jumps over the lazy dog.').



Query Create Function Debugger: fnc_count_vowels2

```

1 DECLARE
2   str text;
3   i integer;
4   len integer;
5   rec record;
6 BEGIN
7   str := upper(p_input);
8   i := 1;
9   len := length(p_input);
10  SELECT 0 AS a, 0 AS e, 0 AS i, 0 AS o, 0 AS u INTO rec;
11 WHILE i <= len LOOP
12   CASE substr(str, i, 1)
13     WHEN 'A' then rec.a := rec.a + 1;
14     WHEN 'E' then rec.e := rec.e + 1;
15     WHEN 'I' then rec.i := rec.i + 1;
16   END CASE;
17   i := i + 1;
18 END LOOP;
19 RETURN rec;
20 END;

```

Cancel Ready

	Parameters	Variables	Result	Messages	Statistics
1	\$1		text	The quick brown fox jumps over the lazy dog.	
2	found		bool	f	
3	str		text	NULL	
4	i		int4	NULL	
5	len		int4	NULL	
6	_Case_Variable_16_		int4	NULL	

Note from the picture above that PostgreSQL created an internal *Case Variable*. Also note that the variable *rec* is not shown in the list of known variables. This is because PostgreSQL still does not know what attributes *rec* will contain. Let's step over some more steps.

Query Create Function Debugger: fnc_count_vowels2

```

5   len integer;
6   rec record;
7 BEGIN
8   str := upper(p_input);
9   i := 1;
10  len := length(p_input);
11  SELECT 0 AS a, 0 AS e, 0 AS i, 0 AS o, 0 AS u INTO rec;
12 WHILE i <= len LOOP
13   CASE substr(str, i, 1)
14     WHEN 'A' then rec.a := rec.a + 1;
15   END CASE;
16   i := i + 1;
17 END LOOP;
18 RETURN rec;
19 END;

```

Cancel Ready

	Parameters	Variables	Result	Messages	Statistics
1	\$1		text	The quick brown fox jumps over the lazy dog.	
2	found		bool	t	
3	str		text	THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.	
4	i		int4	1	
5	len		int4	44	
6	rec	a	int4	0	
7	rec	e	int4	0	
8	rec	i	int4	0	
9	rec	o	int4	0	
10	rec	u	int4	0	
11	_Case_Variable_16_		int4	NULL	

Right after the execution of line 11, *rec* variable comes to life and we can see it has 5 attributes: *a*, *e*, *i*, *o* and *u*, all of the type *int* and having initial value 0.

Now set a breakpoint in line 23 and click the *Resume* button.

Query Create Function Debugger: fnc_count_vowels2

```

17 WHEN 'O' then rec.o := rec.o + 1;
18 WHEN 'U' then rec.u := rec.u + 1;
19 ELSE NULL;
20 END CASE;
21 i := i + 1;
22 END LOOP;
23 RETURN rec.a + rec.e + rec.i + rec.o + rec.u;
24 END;
25

```

	Parameters	Variables	Result	Messages	Statistics
1	\$1		text	The quick brown fox jumps over the lazy dog.	
2	found		bool	t	
3	str		text	THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.	
4	i		int4	45	
5	len		int4	44	
6	rec	a	int4	1	
7	rec	e	int4	3	
8	rec	i	int4	1	
9	rec	o	int4	4	
10	rec	u	int4	2	
11	__Case_Variable_16__		text	NULL	

See how we can inspect every attribute, observing how many of each vowel the text contain. Now let's finish this function.

Query Create Function Debugger: fnc_count_vowels2

```

17 WHEN 'O' then rec.o := rec.o + 1;
18 WHEN 'U' then rec.u := rec.u + 1;
19 ELSE NULL;
20 END CASE;
21 i := i + 1;
22 END LOOP;
23 RETURN rec.a + rec.e + rec.i + rec.o + rec.u;
24 END;
25

```

Finished - Total duration: 1.240 s

	Parameters	Variables	Result	Messages	Statistics
Duration(s)	0.6 0.5 0.4 0.3 0.2 0.1 0	0.0139 0.0133 0.0135 0.0133 0 0.0139 0.0133 0.0132 0.0132 0 0.0136 0.0131	0 0 0 0 0 0 0 0 0 0 0 0	8 9 10 11 12 13 21 13 21 13 15 21	

14 Monitoring Dashboard

OmniDB 2.4.0 introduces a new cool feature called *Monitoring Dashboard*. We know a picture is worth a thousand words, so please take a look:

OMNIDA Connections

Snippets Test

PostgreSQL 9.6.6

- testdb
 - Schemas
 - Extensions
- Databases
- Tablespaces
- Roles
- Replication Slots

Query Monitoring

Refresh All Manage Units

Backends 5 seconds

Backends (max_connections: 100)

Database Size 30 seconds

Database Size (Total)

Size: Top 5 Tables 15 seconds

Size: Top 5 Tables

Activity 15 seconds 2 rows

	datid	datname	pid	usesysid	...
1	23715	testdb	1253	23485	o
2	23715	testdb	1254	23485	o

As you can see, this is a new kind of inner tab showing some charts and grids. This *Monitoring* inner tab is automatically opened once you expand the tree root node (the *PostgreSQL* node). You can keep it open or close it at any time. To open it again, right-click the root node and click on *Dashboard*.

OMNIDA Connections

Snippets Test

PostgreSQL 9.6.6

- testdb
 - Sc
 - Ex
 - Doc: PostgreSQL
 - Doc: SQL Language
 - Tables
 - Replication Slots

Query

1

Refresh

Monitoring

- Dashboard
- Backends
- Replication

Data Messages Explain

The dashboard is composed of handy information rectangles called *Monitoring Units*. Here is an example of Monitoring Unit and its interface elements:



- 1: Title of the Monitoring Unit;
- 2: Refresh the Monitoring Unit. Depending on the type, clicking on this button will refresh the entire drawing area or just make the chart acquire a new set of values;
- 3: Pause the Monitoring Unit;
- 4: Interval in seconds for automatic refreshing;
- 5: Remove the Monitoring Unit of the Monitoring Dashboard;
- 6: Drawing area, that will be different depending on the type of the Monitoring Unit.

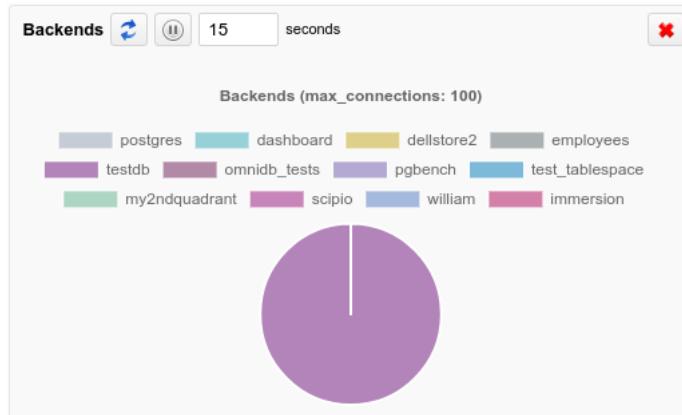
14.0.0.1 Types of Monitoring Units

Currently there are 3 types of Monitoring Units:

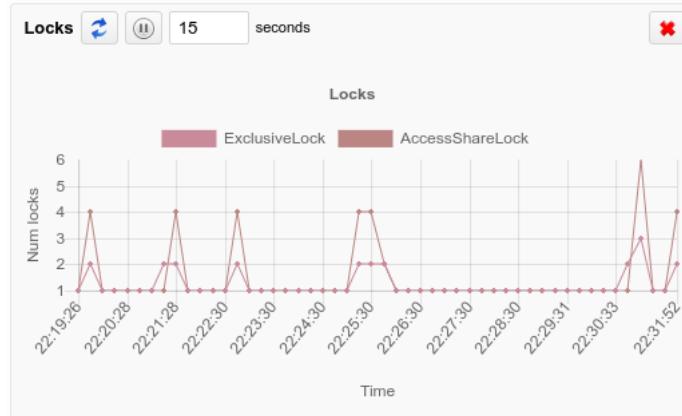
- **Grid:** The most simple kind, just executes a query from time to time and shows the results in a data grid.

	datid	datname	pid	usesysid	username	application_name	client_addr
1	23715	testdb	3217	23485	omnidb		127.0.0.1

- **Chart:** Every time it refreshes, it renders a new complete chart. The old set of values is lost. This is most useful for pie charts, but other kind of charts can be used too.

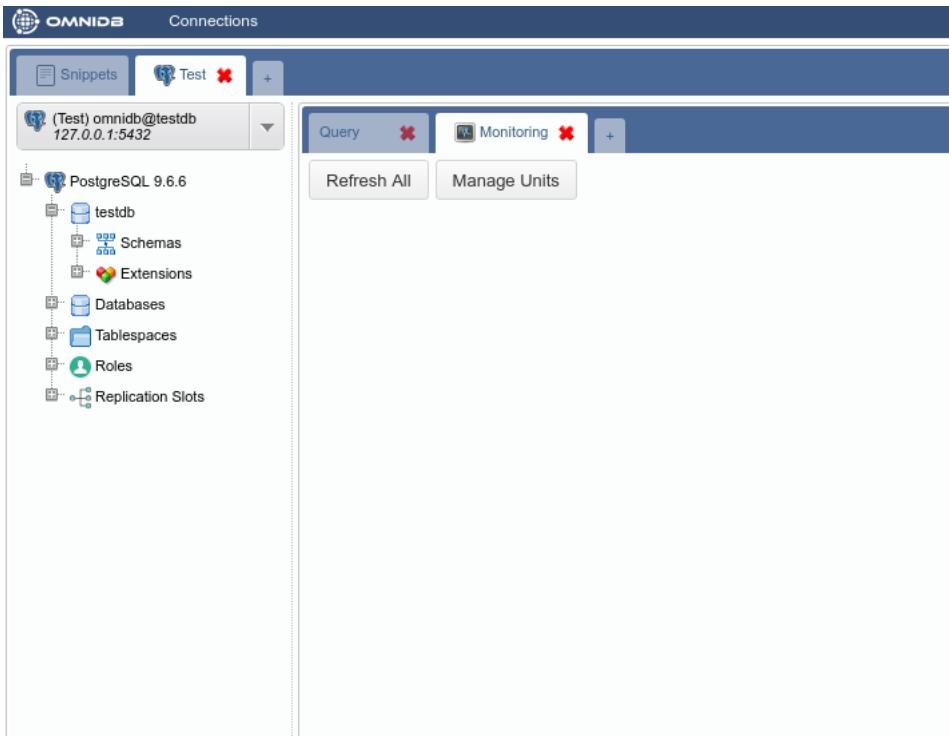


- **Chart-Append:** Perhaps this is the most useful kind of Monitoring Unit. It is a chart that appends a new set of values every time it refreshes. Line or bar charts fit best for this type. The last 50 set of values are kept by the component client-side to be viewed by the user.

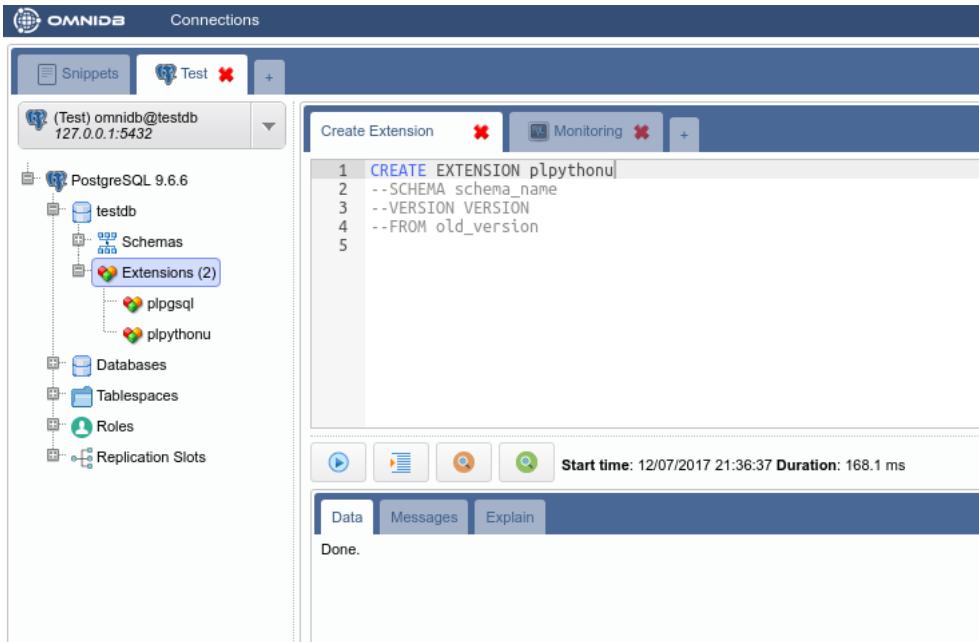


14.0.0.2 Showing and hiding units in the dashboard

If you click in the button *Refresh All*, then all Monitoring Units will be refreshed at once. You can also remove undesired Monitoring Units by clicking in the *Remove* button. Let us go ahead and remove all units from the dashboard, making it empty:



All Monitoring Units that come with OmniDB are open source and available in this [repository](#) (feel free to contribute). But be aware that some Monitoring Units require the `plpythonu` script to be installed in the database. Please refer to the instructions specific to your operating system on how to install `plpythonu` if you desire to use and create Monitoring Units that use `plpythonu`.



Now that our dashboard is empty, let us add some units. Click on the *Manage Units* button.

	Actions	Title	Type
1	<input checked="" type="checkbox"/>	Backends	Chart
2	<input checked="" type="checkbox"/>	Database Size	Chart
3	<input checked="" type="checkbox"/>	Backends	Chart (Append)
4	<input checked="" type="checkbox"/>	Bloat: Top 5 Tables	Chart (Append)
5	<input checked="" type="checkbox"/>	CPU Usage	Chart (Append)
6	<input checked="" type="checkbox"/>	Database Size	Chart (Append)
7	<input checked="" type="checkbox"/>	Locks	Chart (Append)
8	<input checked="" type="checkbox"/>	Master: Replication Lag	Chart (Append)
9	<input checked="" type="checkbox"/>	Memory Usage	Chart (Append)
10	<input checked="" type="checkbox"/>	Size: Top 5 Tables	Chart (Append)

Click on the green action to add the Monitoring Units called *CPU Usage* and *Memory Usage*. Bear in mind that both units require *plythonu* extension in the database. *CPU Usage* also requires that the tool *mpstat* should be installed in the server. Also both units are of type *Chart-Append*. Wait for some seconds and you will have a dashboard like this:

Memory Usage 5 seconds

System Memory Usage (Total: 7888MB)

Time

CPU Usage 10 seconds

CPU Usag

Value

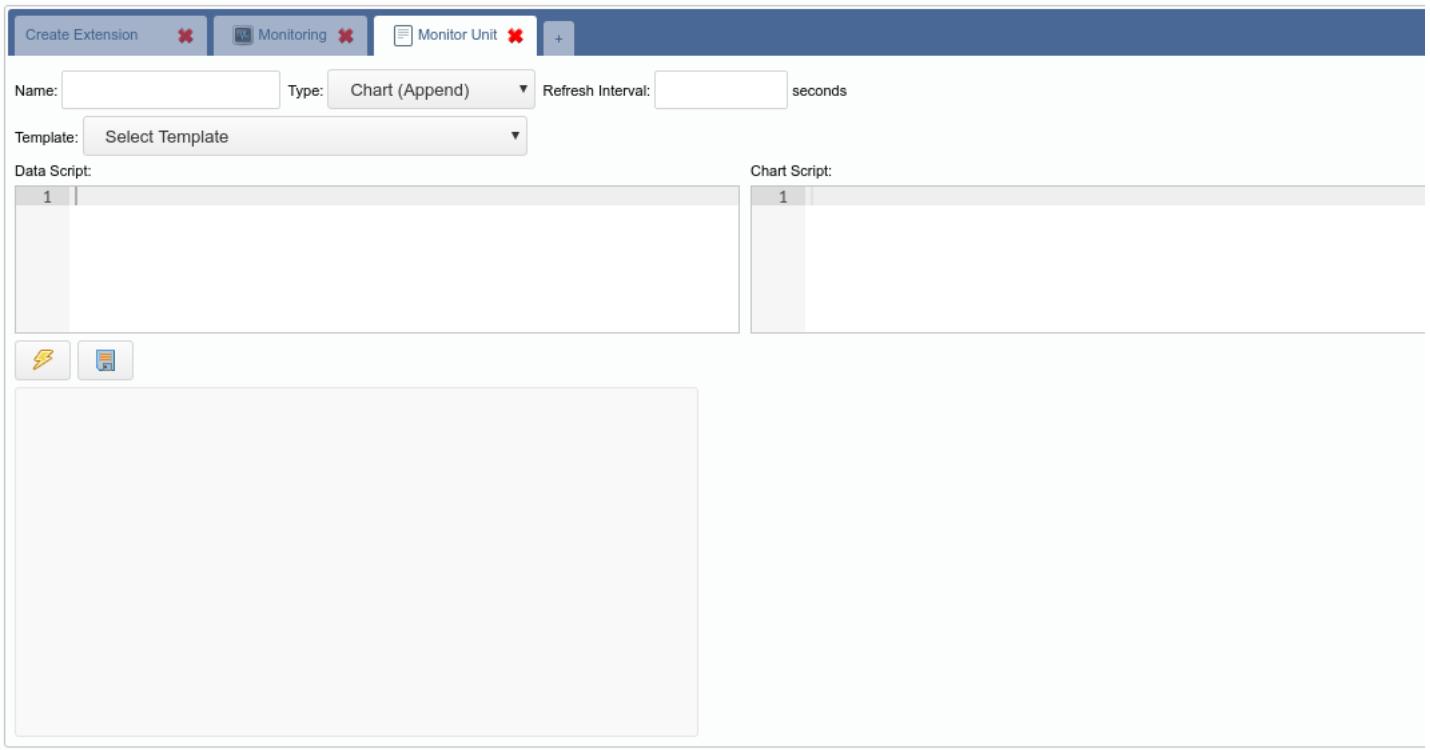
Time

In a similar way, you can add and remove any unit you want to customize the dashboard the way you want.

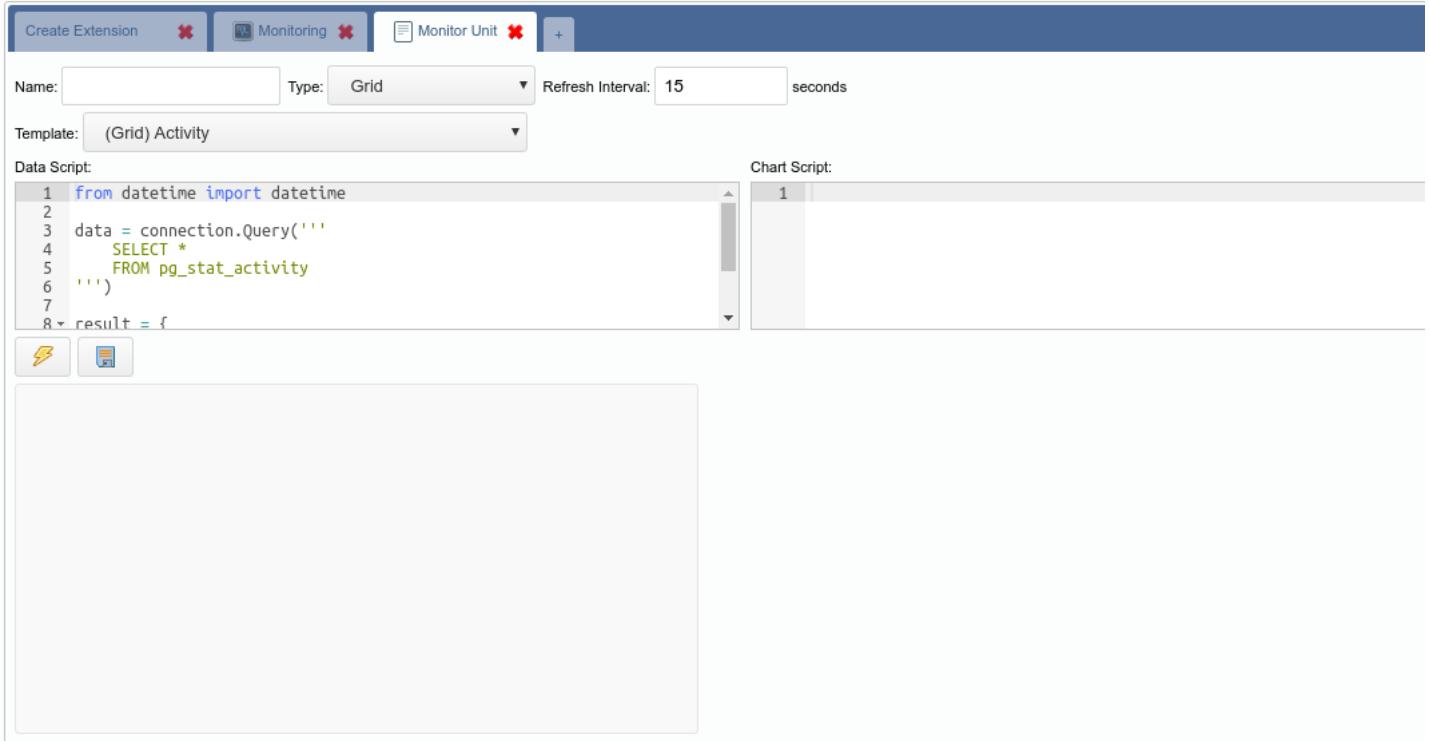
14.0.0.3 Writing custom Monitoring Units: Grid

OmniDB provides you the power to write your own units and customize existing ones. Everything is done through Python scripts that run inside a sandbox. Beware that units powered by *plythonu* can have access to the file system the database user also has access to, and any Monitoring Unit have the same permission as the database user you configured in the Connection.

To create a new Monitoring Unit, click on the *Manage Units* button in the dashboard, then click on the *New Unit* button. It will open a new kind of inner tab like this:



The easiest way to write a custom unit is to use an existing one as template. Go ahead and select the *(Grid) Activity* template:



Note how OmniDB fills the *Data Script* source code. This script is responsible for generating data for the unit every time it refreshes. As a grid unit is nothing else but a grid of data, we can rely on only this script for now.

Now let us take a look at the source code of this template:

```
from datetime import datetime

data = connection.Query('''
    SELECT *
    FROM pg_stat_activity
''')

result = {
    "columns": data.Columns,
    "data": data.Rows
}
```

It is simple enough. It executes an SQL query into the current connection using the reserved `connection` variable. Also, the grid unit type expects its results in a JSON variable that must be called `result` and must have the attributes "`columns`" (an array of column names) and "`data`" (an array of rows, each row being an array of values). The `connection.Query()` function already does the job pretty well, so let us just change the SQL query this way:

```

from datetime import datetime

data = connection.Query('''
    SELECT random() as "Random Number"
''')

result = {
    "columns": data.Columns,
    "data": data.Rows
}

```

Copy and paste the above Python code into the *Data Script* text field and then click on the *Test* (lightning) button:

The screenshot shows the 'Monitor Unit' configuration screen. In the 'Data Script' section, the following Python code is pasted:

```

3 data = connection.Query('''
4     SELECT random() as "Random Number"
5 ''')
6
7 result = []
8     "columns": data.Columns,
9     "data": data.Rows
10 }

```

Below the script, there is a preview area showing a single row of data:

	Random Number
1	0.749598043505102

Note how the grid was rendered in the preview drawing area. You can click the *Test* button as many times as you want. Now we will give the unit a name, set refresh interval and then hit the *Save* button:

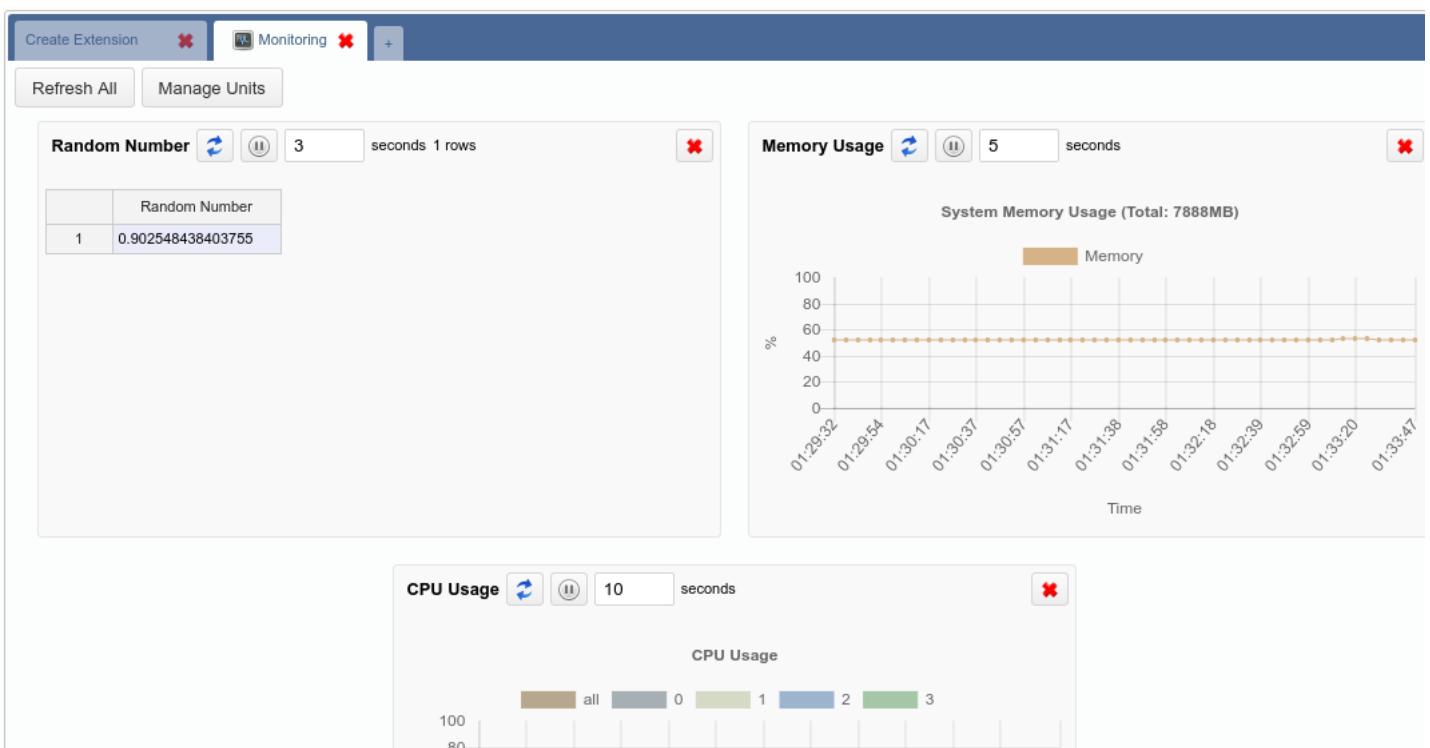
The screenshot shows the 'Monitor Unit' configuration screen with the following settings:

- Name: Random Number
- Type: Grid
- Refresh Interval: 3 seconds

The 'Data Script' section contains the same Python code as before.

A modal dialog box is displayed in the center of the screen with the message "Monitor unit saved." and an "Ok" button.

Click the *OK* button and then close the edit tab. Our new Monitoring Unit will be in the list of available units. As we created this unit, we can either add it to the dashboard, edit it or remove it. Let us add it to the dashboard (green action):



14.0.0.4 Writing custom Monitoring Units: Chart

Click in the *Manage Units* button and then in the *New Unit* button. This time we will create a Chart Monitoring Unit. So choose *(Chart) Database Size* as a template.

The screenshot shows the pgAdmin Monitor Unit creation dialog for a new unit named "DatabaseSize". The "Type" is set to "Chart (No Append)". The "Refresh Interval" is set to 30 seconds. The "Template" is set to "(Chart) Database Size".

Data Script:

```

1 from datetime import datetime
2 from random import randint
3 |
4 databases = connection.Query('''
5     SELECT d.datname AS datname,
6         round(pg_catalog.pg_database_size(d.datname)/1048576.0,2)
7     FROM pg_catalog.pg_database d
8 ''')

```

Chart Script:

```

1 total_size = connection.ExecuteScalar('''
2     SELECT round(sum(pg_catalog.pg_database_size(datname)/1048576
3         FROM pg_catalog.pg_database
4         WHERE NOT datistemplate
5     ''')
6 |
7     result = {
8

```

The source code of this kind of unit is more complex. There are two scripts:

- **Data Script**: Executed every time the unit is refreshed;
- **Chart Script**: Executed only at the beginning to build the chart.

Let us take a look at the **Data Script**:

```

from datetime import datetime
from random import randint

databases = connection.Query('''
    SELECT d.datname AS datname,
        round(pg_catalog.pg_database_size(d.datname)/1048576.0,2) AS size
    FROM pg_catalog.pg_database d
    WHERE d.datname not in ('template0','template1')
''')

data = []
color = []
label = []

```

```

for db in databases.Rows:
    data.append(db["size"])
    color.append("rgb(" + str(randint(125, 225)) + "," + str(randint(125, 225)) + "," + str(randint(125, 225)) + ")")
    label.append(db["datname"])

result = {
    "labels": label,
    "datasets": [
        {
            "data": data,
            "backgroundColor": color,
            "label": "Dataset 1"
        }
    ]
}

```

Here we can see that the reserved variable `connection` is still being used to retrieve data from the database. Bear in mind that this variable is always pointing to the current Connection. Three arrays are needed to build a set of values:

- `data`: One value per dataset. Each dataset will be a chart dataset (a line in a line chart, for example);
- `color`: One color per dataset;
- `label`: One label per dataset.

This way, `data[0]`, `color[0]` and `label[0]` refer to the first dataset, while `data[1]`, `color[1]` and `label[1]` refer to the second dataset, and so on.

This script must return a variable called `result` and also needs to be a JSON like in the above script.

So right now you are probably guessing that you just need to change the SQL query to make the chart behave different. Well, in terms of data and datasets, you guessed right. So let's change the SQL query of this chart to compare sizes of tables of schema `public`. Also change the references from `datname` to `tablename`, as we have changed the column name.

```

from datetime import datetime
from random import randint

databases = connection.Query('''
    SELECT c.relname as tablename,
           round(pg_catalog.pg_total_relation_size(c.oid)/1048576.0,2) AS size
    FROM pg_catalog.pg_class c
    INNER JOIN pg_catalog.pg_namespace n
    ON n.oid = c.relnamespace
    WHERE n.nspname = 'public'
      AND c.relkind = 'r'
''')

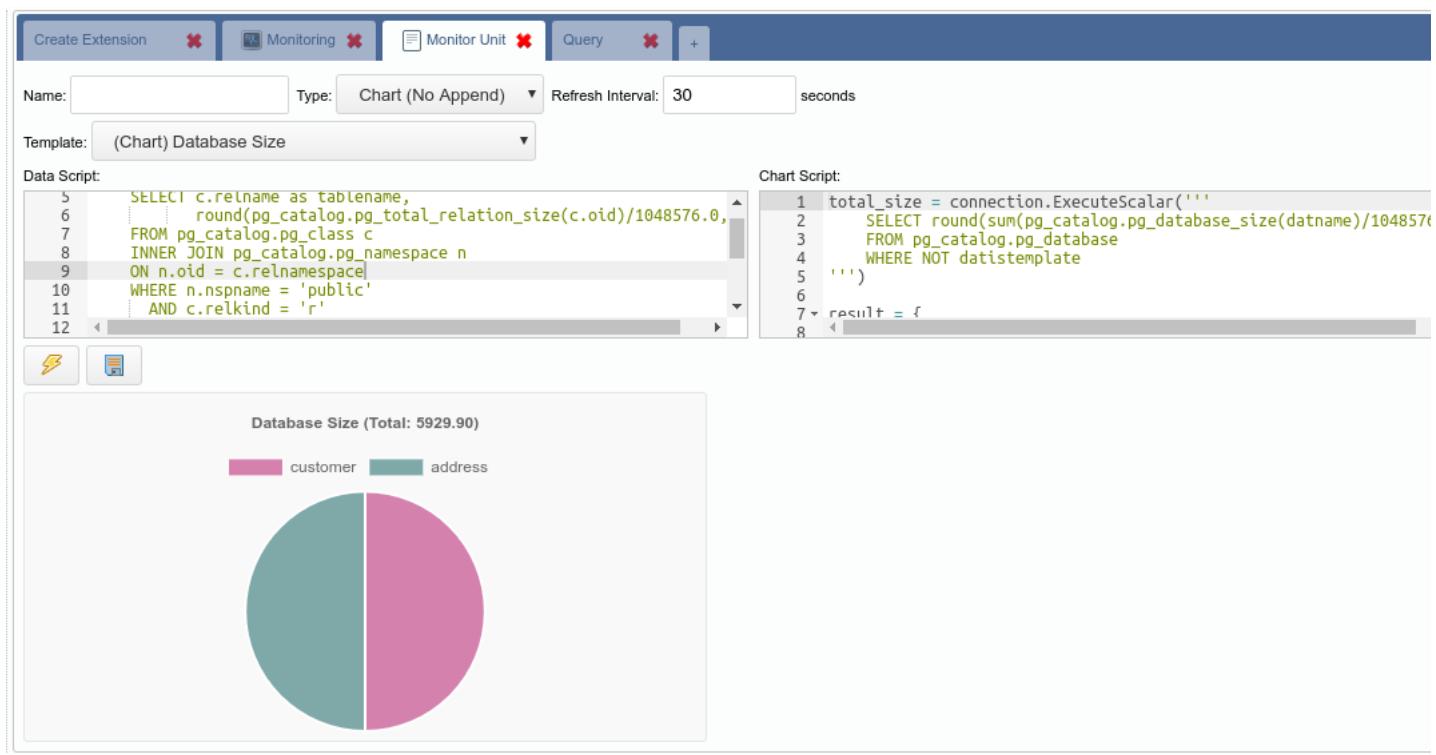
data = []
color = []
label = []

for db in databases.Rows:
    data.append(db["size"])
    color.append("rgb(" + str(randint(125, 225)) + "," + str(randint(125, 225)) + "," + str(randint(125, 225)) + ")")
    label.append(db["tablename"])

result = {
    "labels": label,
    "datasets": [
        {
            "data": data,
            "backgroundColor": color,
            "label": "Dataset 1"
        }
    ]
}

```

Copy and paste the above script into the *Data Script* field and then hit the *Test* button:



Apparently the chart is almost done. We need to fix the title, it still says *Database Size*, when this chart is about table size. Any information about the format of the chart itself is defined in the *Chart Script* text field. Let us understand the current source code:

```

total_size = connection.ExecuteScalar('''
SELECT round(sum(pg_catalog.pg_database_size(datname)/1048576.0),2)
FROM pg_catalog.pg_database
WHERE NOT datistemplate
''')

result = {
  "type": "pie",
  "data": None,
  "options": {
    "responsive": True,
    "title":{
      "display":True,
      "text":"Database Size (Total: " + str(total_size) + ")"
    }
}
}

```

Easy enough. We can make use of the reserved variable *connection* to retrieve data in the *Chart Script* too. This is mainly used to put information in the chart title. The variable *result* must be defined here. Note how its JSON value defines a pie chart and the title. So we just need to change the query and adjust the title, this way:

```

total_size = connection.ExecuteScalar('''
SELECT round(sum(pg_catalog.pg_total_relation_size(c.oid)/1048576.0),2) AS size
FROM pg_catalog.pg_class c
INNER JOIN pg_catalog.pg_namespace n
ON n.oid = c.relnamespace
WHERE n.nspname = 'public'
AND c.relkind = 'r'
''')

result = {
  "type": "pie",
  "data": None,
  "options": {
    "responsive": True,
    "title":{
      "display":True,
      "text":"Table Size (Total: " + str(total_size) + ")"
    }
}
}

```

Copy and paste the above Python code into the *Chart Script*. Then click in the *Test* button:

Create Extension X Monitoring X Monitor Unit X Query X +

Name: Type: Chart (No Append) Refresh Interval: 30 seconds

Template: (Chart) Database Size

Data Script:

```

5   SELECT c.relname as tablename,
6   ... round(pg_catalog.pg_total_relation_size(c.oid)/1048576.0,
7   ... FROM pg_catalog.pg_class c
8   ... INNER JOIN pg_catalog.pg_namespace n
9   ... ON n.oid = c.relnamespace
10  ... WHERE n.nspname = 'public'
11  ... AND c.relkid = 'r'
12

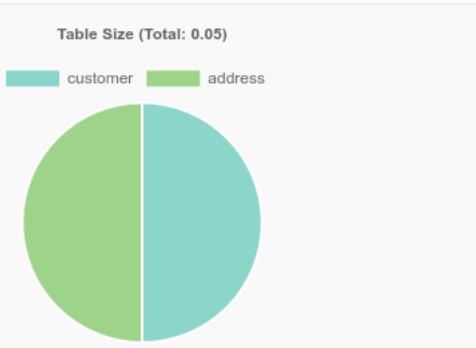
```

Chart Script:

```

14   ...
15   ...
16   ...
17   ...
18   ...
19   ...
20   ...

```



Now that the chart finally works the way we want, we can give it a title, adjust the refresh interval and then click in the *Save* button. After that we can add it to the dashboard.

Create Extension X Monitoring X Monitor Unit X Query X +

Name: Table Size Type: Chart (No Append) Refresh Interval: 5 seconds

Template: (Chart) Database Size

Data Script:

```

5   SELECT c.relname as tablename,
6   ... round(pg_catalog.pg_total_relation_size(c.oid)/1048576.0,
7   ... FROM pg_catalog.pg_class c
8   ... INNER JOIN pg_catalog.pg_namespace n
9   ... ON n.oid = c.
10  ... WHERE n.nspna
11  ... AND c.relkid
12

```

Chart Script:

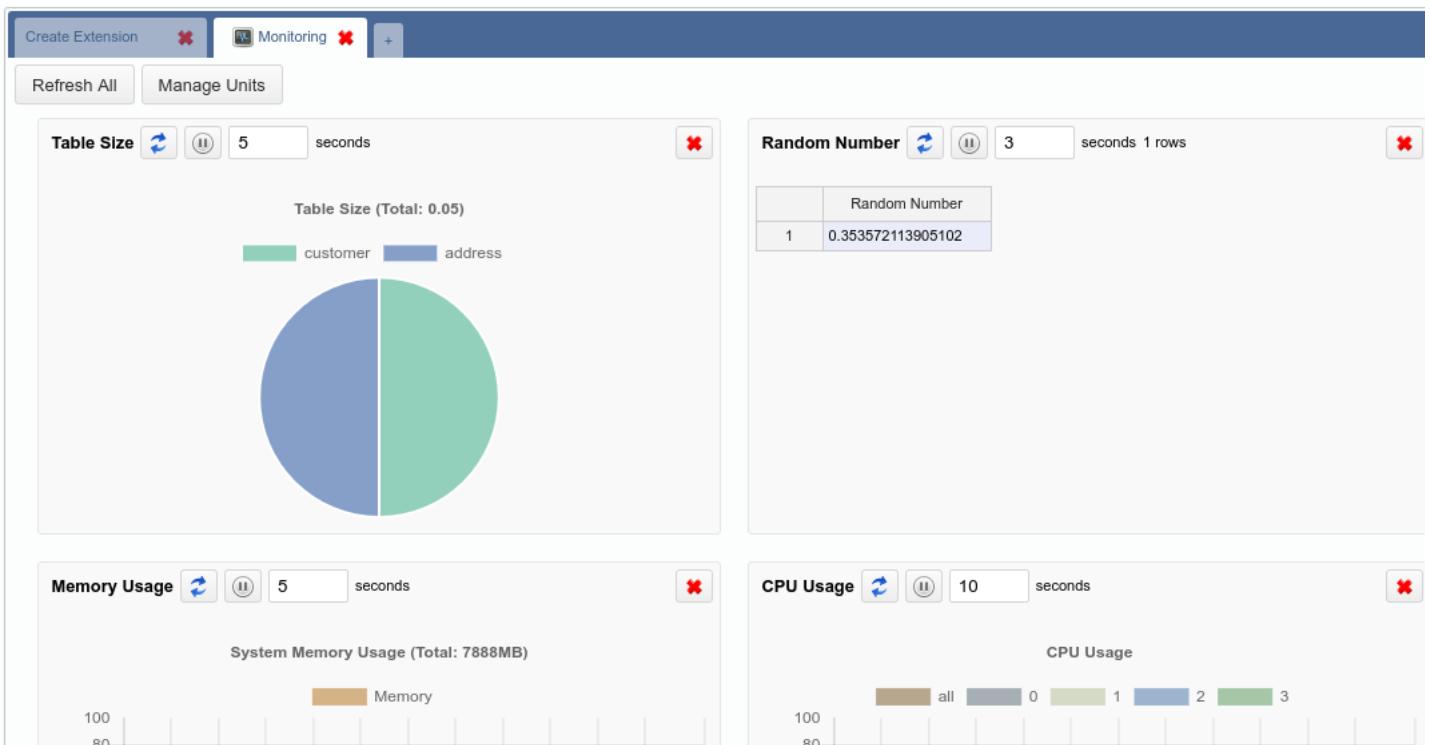
```

14   ...
15   ...
16   ...
17   ...

```

Monitor unit saved.





14.0.0.5 Writing custom Monitoring Units: Chart-Append

Now for the last, but most interesting kind of Monitoring Unit: *Chart-Append*. It is interesting because there is a wide range of applications for these units, since they keep recent historic data that allows us to see a comparison of values.

Go ahead and add a new chart using (*Chart (Append)*) *Size: Top 5 Tables* as template:

The screenshot shows the configuration for a new Monitor Unit:

- Name:** [empty input field]
- Type:** Chart (Append)
- Refresh Interval:** 15 seconds
- Template:** (Chart (Append)) Size: Top 5 Tables
- Data Script:**

```

1 from datetime import datetime
2 from random import randint
3
4 tables = connection.Query('''
5     SELECT nspname || '.' || relname AS relation,
6         round(pg_relation_size(c.oid)/1048576.0,2) AS size
7     FROM pg_class c
8 ''')
    
```
- Chart Script:**

```

1 result = {
2     "type": "line",
3     "data": None,
4     "options": {
5         "responsive": True,
6         "title": {
7             "display": True,
8             "text": "Size: Top 5 Tables"
9         }
10    }
11 }
    
```

Now take a look at the source code of both *Data Script* and *Chart Script*. It is not too different from the Chart units. The dataset creation is a bit more complex as it involves other JSON settings, but that's all.

As an exercise, based on this chart, create another one called *Size: Top 20 Tables*. It should look like this:

Create Extension Monitoring Monitor Unit

Name: Size: Top 20 Tables Type: Chart (Append) Refresh Interval: 15 seconds

Template: (Chart (Append)) Size: Top 5 Tables

Data Script:

```

27     "data": [table["size"]]
28   })
29
30  result = {
31    "labels": [datetime.now().strftime('%H:%M:%S')],
32    "datasets": datasets
33
34 }
35
36
  
```

Chart Script:

```

29
30  "display": True,
31  "labelString": "Size (MB)"
32
33
34
35
36
  
```

Size: Top 20 Tables

Size (MB)

Time

Legend:

- omnidb.statistics
- omnidb.variables
- omnidb.contexts
- public.customer
- public.address
- omnidb.bar
- public.address_add_id_seq
- public.customer_cust_id_seq
- omnidb.foo

Now save it and add it to your dashboard:

Monitoring

Refresh All Manage Units

Size: Top 20 Tables 15 seconds

Table Size 5 seconds

Random Number 3 seconds 1 rows

Memory Usage 5 seconds

Size: Top 20 Tables

Size (MB)

Time

Legend:

- omnidb.statistics
- omnidb.variables
- omnidb.contexts
- public.customer
- public.address
- omnidb.bar
- public.address_add_id_seq
- public.customer_cust_id_seq
- omnidb.foo

Table Size (Total: 0.05)

customer address

System Memory Usage (Total: 7888MB)

Memory